


 DeepLearning.AI

















Stanford  
ONLINE



# Decision Trees

## Decision Tree Model

# Cat classification example

	Ear shape ( $x_1$ )	Face shape ( $x_2$ )	Whiskers ( $x_3$ )	Cat
	Pointy 	Round 	Present 	1
	Floppy 	Not round 	Present	1
	Floppy	Round	Absent 	0
	Pointy	Not round	Present	0
	Pointy	Round	Present	1
	Pointy	Round	Absent	1
	Floppy	Not round	Absent	0
	Pointy	Round	Absent	1
	Floppy	Round	Absent	0
	Floppy	Round	Absent	0

Categorical (discrete values)

$X$

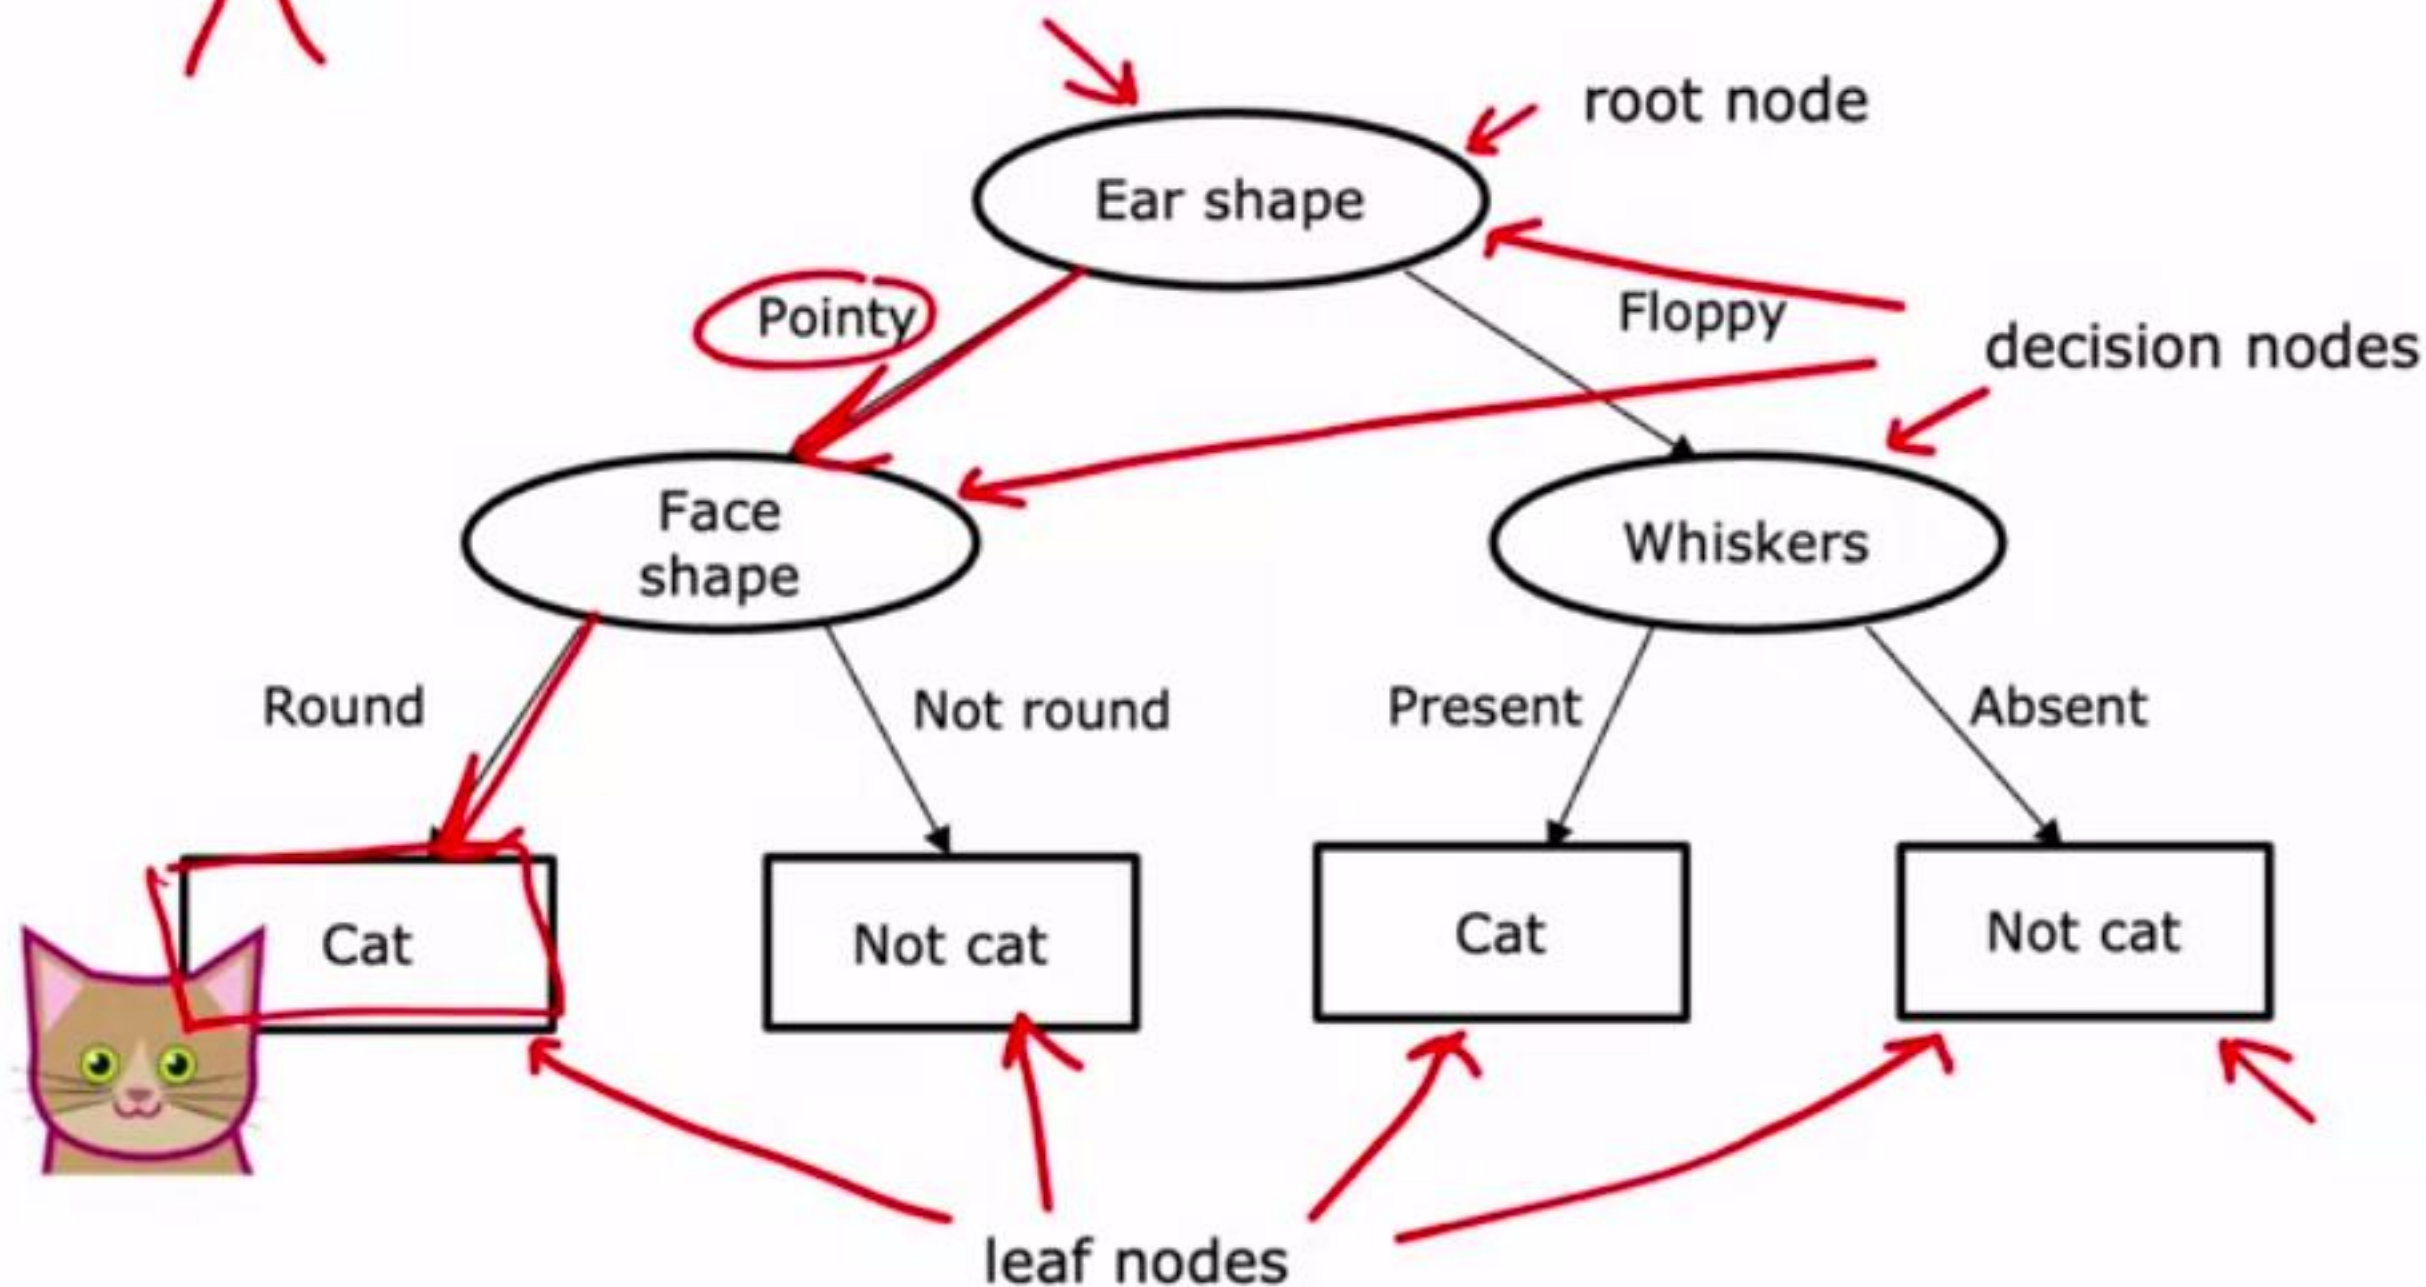
$y$





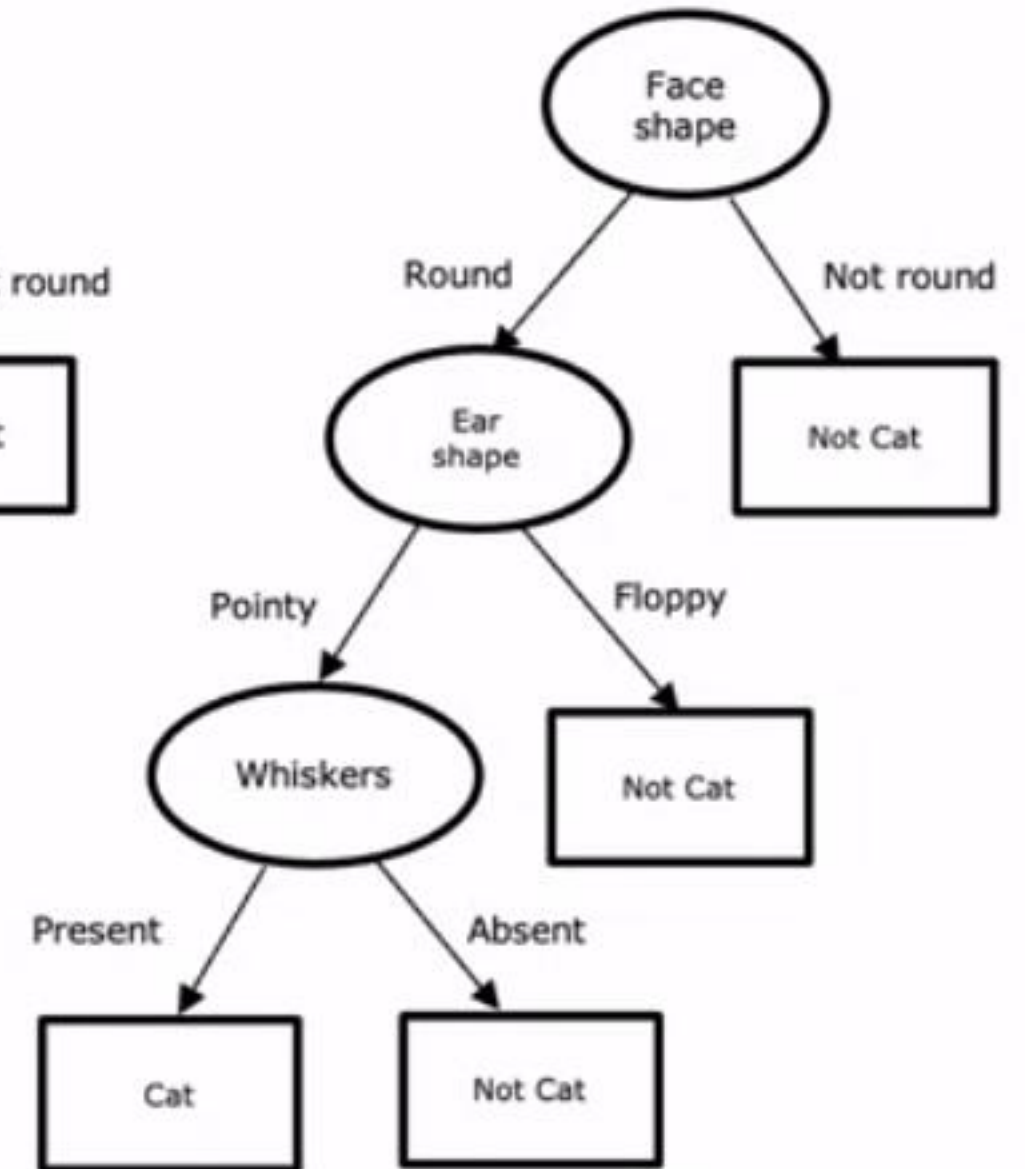

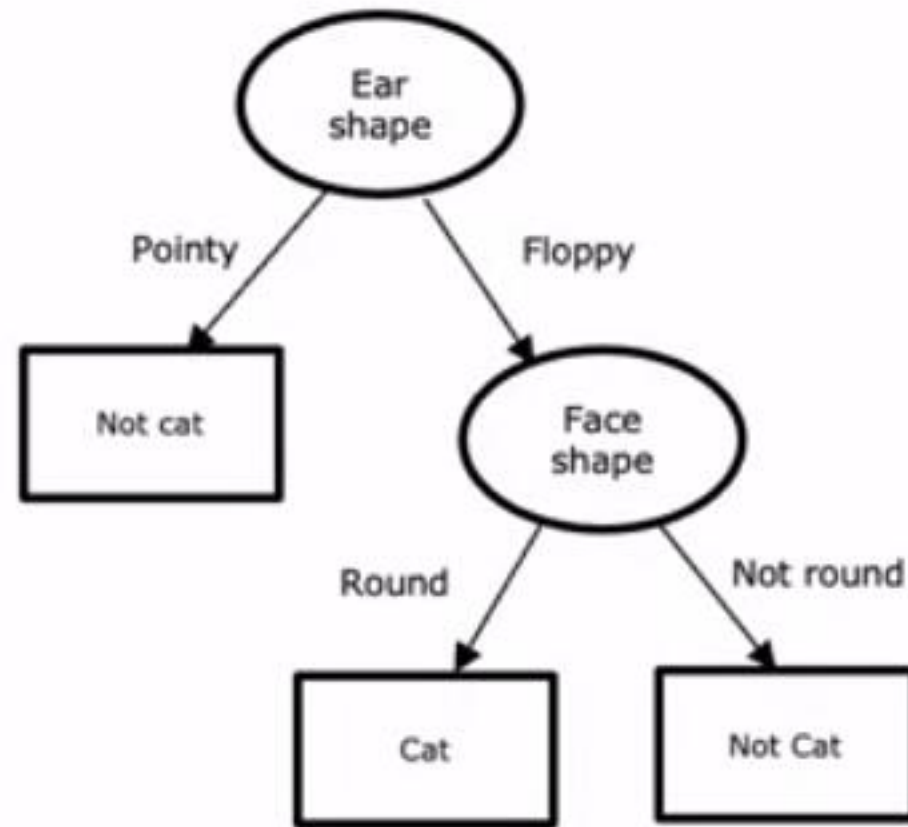
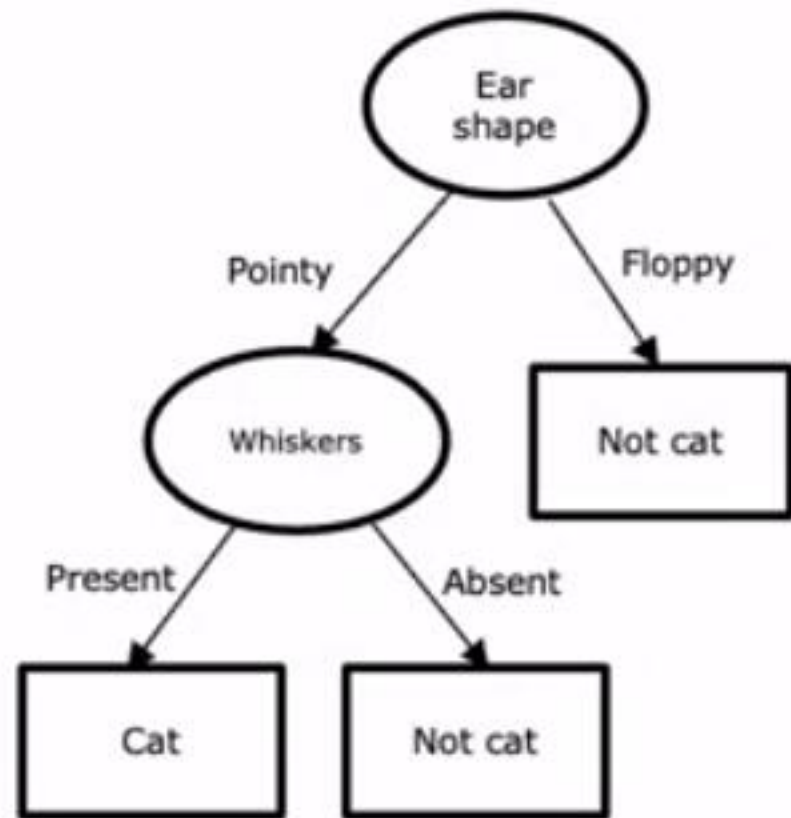
# Decision Tree

New test example




Ear shape: Pointy  
Face shape: Round  
Whiskers: Present

# Decision Tree



 DeepLearning.AI

Stanford  
ONLINE

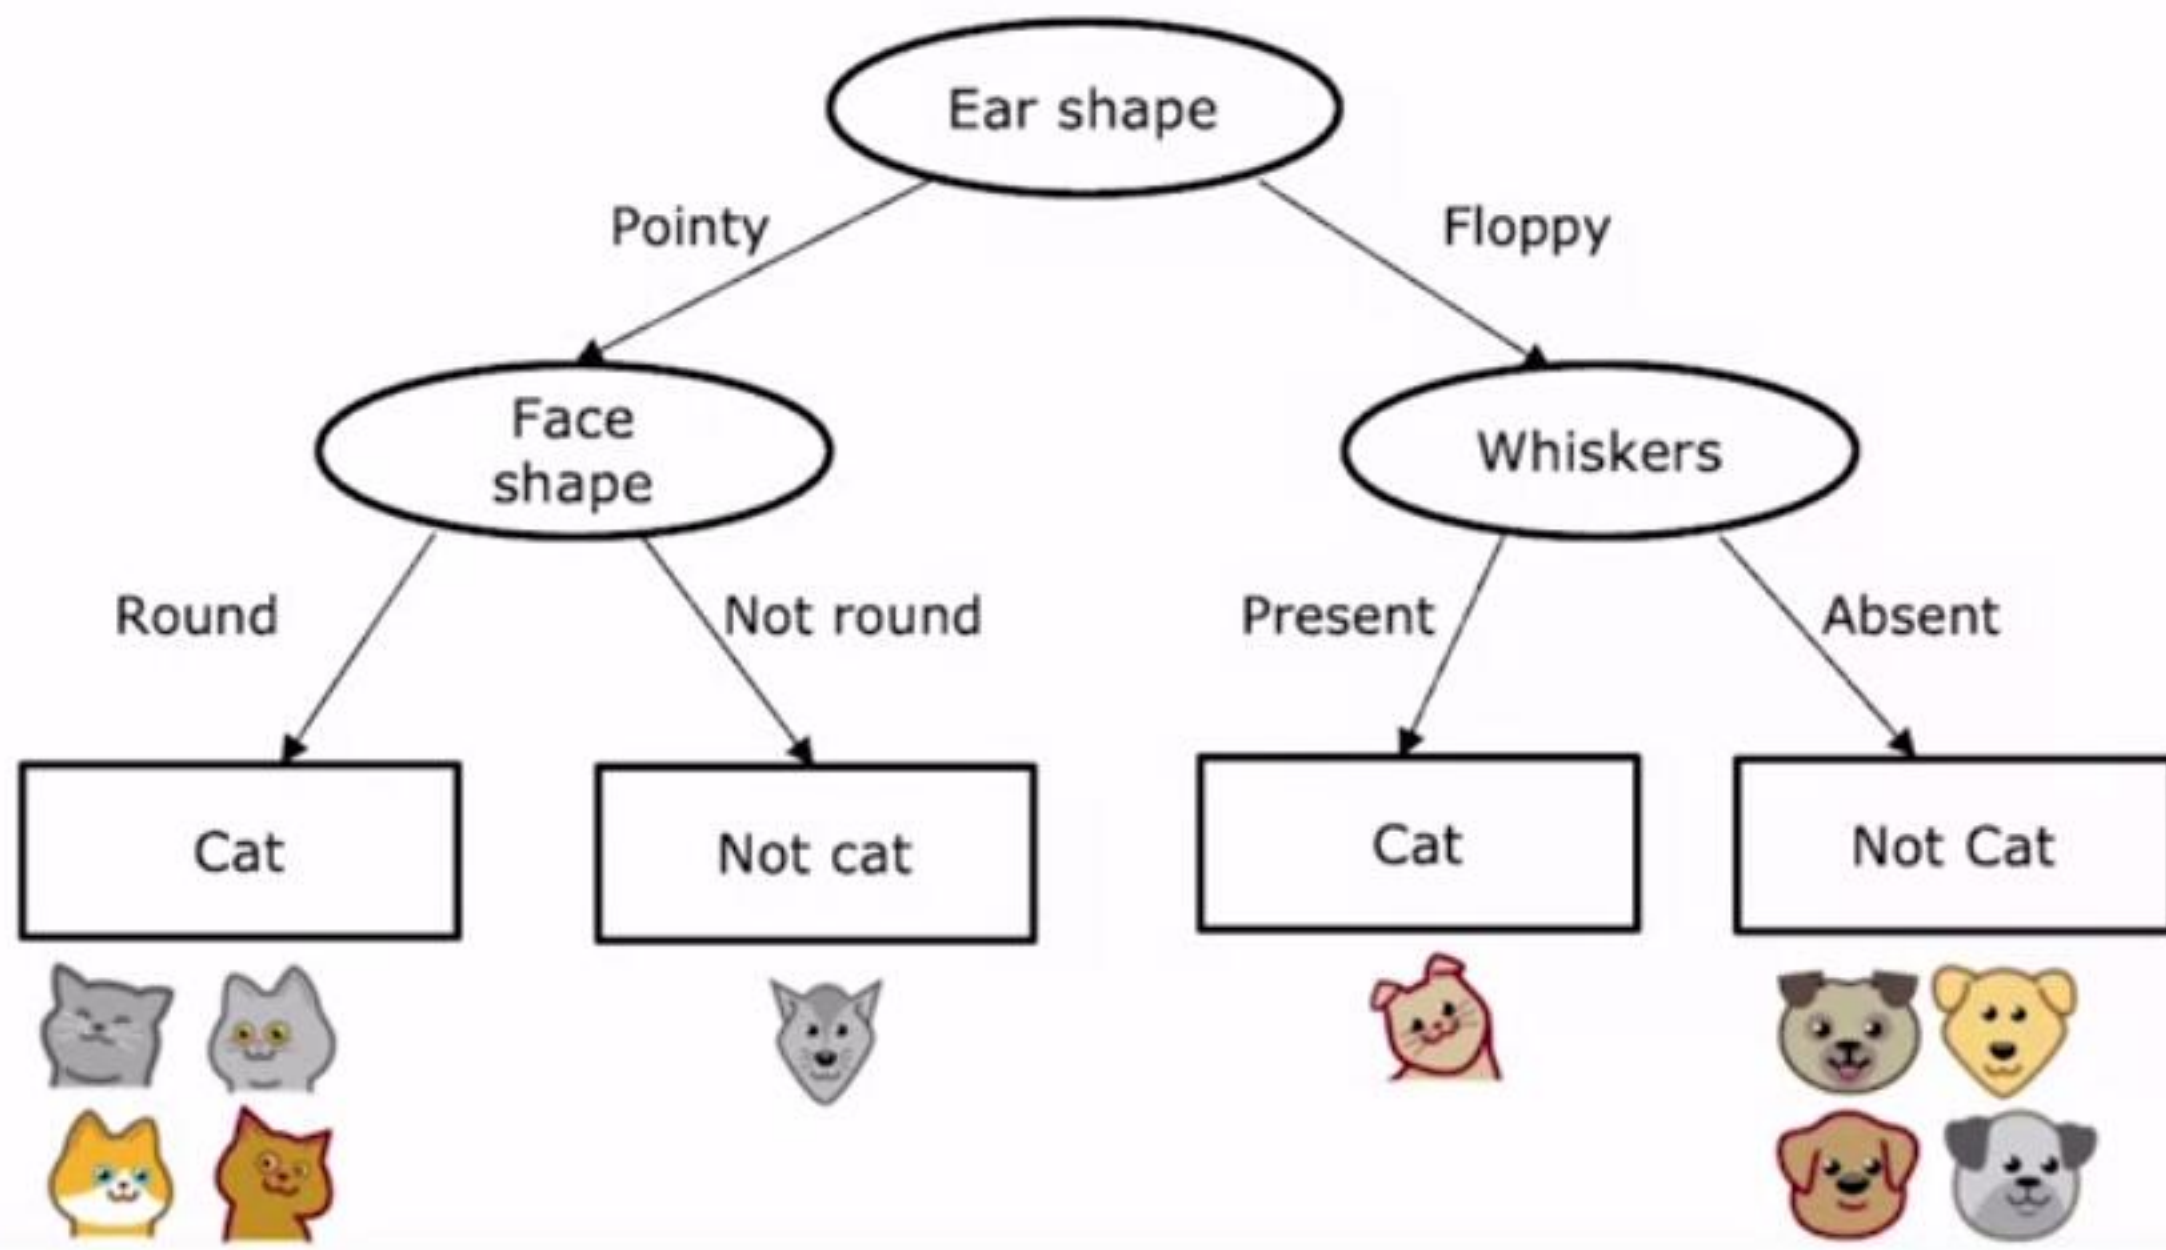


# Decision Trees

## Learning Process




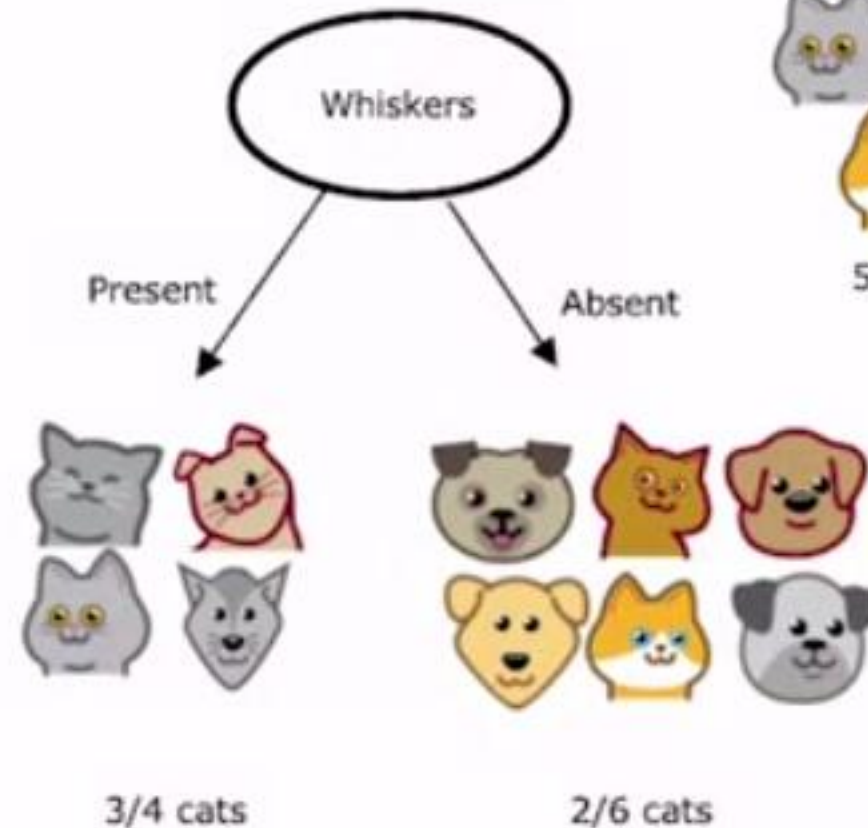


# Decision Tree Learning



# Decision Tree Learning

**Decision 1:** How to choose what feature to split on at each node?

Maximize purity (or minimize impurity)






# Decision Tree Learning

**Decision 2:** When do you stop splitting?


- When a node is 100% one class
- When splitting a node will result in the tree exceeding a maximum depth
- When improvements in purity score are below a threshold
- When number of examples in a node is below a threshold





 DeepLearning.AI

Stanford  
ONLINE




# Decision Tree Learning

## Measuring purity

# Entropy as a measure of impurity

$p_1$  = fraction of examples that are cats



					
Dog	Dog	Dog	Dog	Dog	Dog
					
Cat	Cat	Dog	Dog	Dog	Dog
					
Cat	Cat	Cat	Dog	Dog	Dog
					
Cat	Cat	Cat	Cat	Cat	Dog
					
Cat	Cat	Cat	Cat	Cat	Cat

$$p_1 = 0 \quad H(p_1) = 0$$

$$p_1 = 2/6 \quad H(p_1) = 0.92$$


$$p_1 = 3/6 \quad H(p_1) = 1$$

$$p_1 = 5/6 \quad H(p_1) = 0.65$$

$$p_1 = 6/6 \quad H(p_1) = 0$$

# Entropy as a measure of impurity

$p_1$  = fraction of examples that are cats



$$p_0 = 1 - p_1$$


$$\begin{aligned} H(p_1) &= -p_1 \log_2(p_1) - p_0 \log_2(p_0) \\ &= -p_1 \log_2(p_1) - (1 - p_1) \log_2(1 - p_1) \end{aligned}$$

Note: " $0 \log(0)$ " = 0





Stanford  
ONLINE




# Decision Tree Learning

Choosing a split: Information  
Gain

# Choosing a split

$$p_1 = 5/10 = 0.5$$

$$H(0.5) = 1$$




$$p_1 = 4/5 = 0.8 \quad p_1 = 1/5 = 0.2$$

$$H(0.8) = 0.72 \quad H(0.2) = 0.72$$

$$H(0.5) - \left( \frac{5}{10} H(0.8) + \frac{5}{10} H(0.2) \right)$$

$$= 0.28$$

$$H(0.5) = 1$$



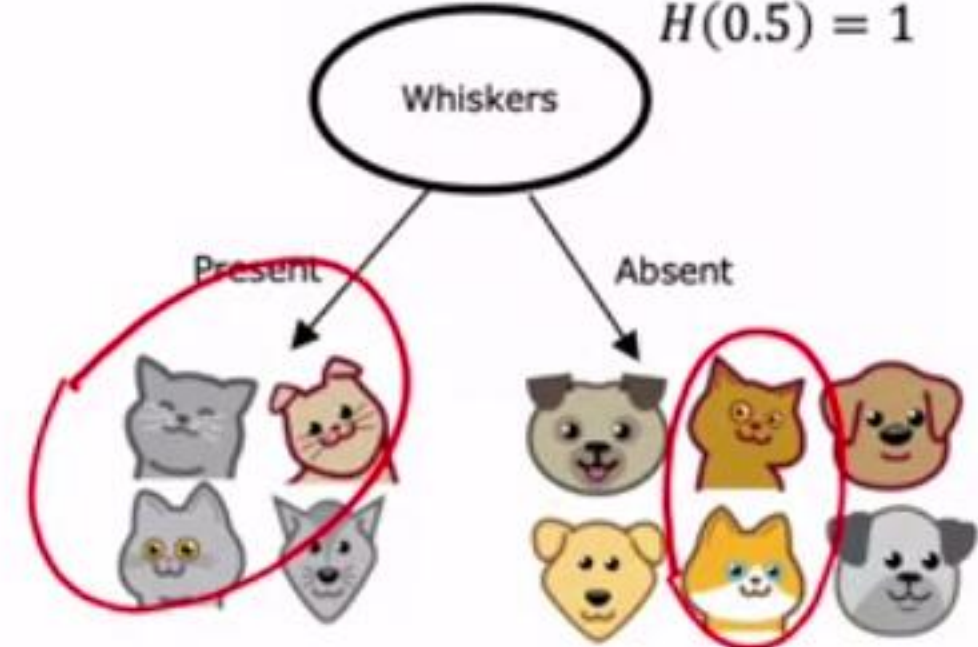
$$p_1 = 4/7 = 0.57 \quad p_1 = 1/3 = 0.33$$

$$H(0.57) = 0.99 \quad H(0.33) = 0.92$$

$$H(0.5) - \left( \frac{7}{10} H(0.57) + \frac{3}{10} H(0.33) \right)$$

$$= 0.03$$

$$H(0.5) = 1$$



$$p_1 = 3/4 = 0.75 \quad p_1 = 2/6 = 0.33$$

$$H(0.75) = 0.81 \quad H(0.33) = 0.92$$

$$H(0.5) - \left( \frac{4}{10} H(0.75) + \frac{6}{10} H(0.33) \right)$$

$$= 0.12$$


Information gain



# Information Gain



$$p_1^{\text{root}} = 5/10 = 0.5$$



$$p_1^{\text{left}} = 4/5$$

$$w^{\text{left}} = 5/10$$

$$p_1^{\text{right}} = 1/5$$

$$w^{\text{right}} = 5/10$$


Information gain

$$= H(p_1^{\text{root}}) - \left( w^{\text{left}} H(p_1^{\text{left}}) + w^{\text{right}} H(p_1^{\text{right}}) \right)$$





Stanford  
ONLINE




# Decision Tree Learning

## Putting it together

# Decision Tree Learning


- Start with all examples at the root node
- Calculate information gain for all possible features, and pick the one with the highest information gain
- Split dataset according to selected feature, and create left and right branches of the tree
- Keep repeating splitting process until stopping criteria is met:
  - When a node is 100% one class
  - When splitting a node will result in the tree exceeding a maximum depth
  - Information gain from additional splits is less than threshold
  - When number of examples in a node is below a threshold

# Recursive splitting

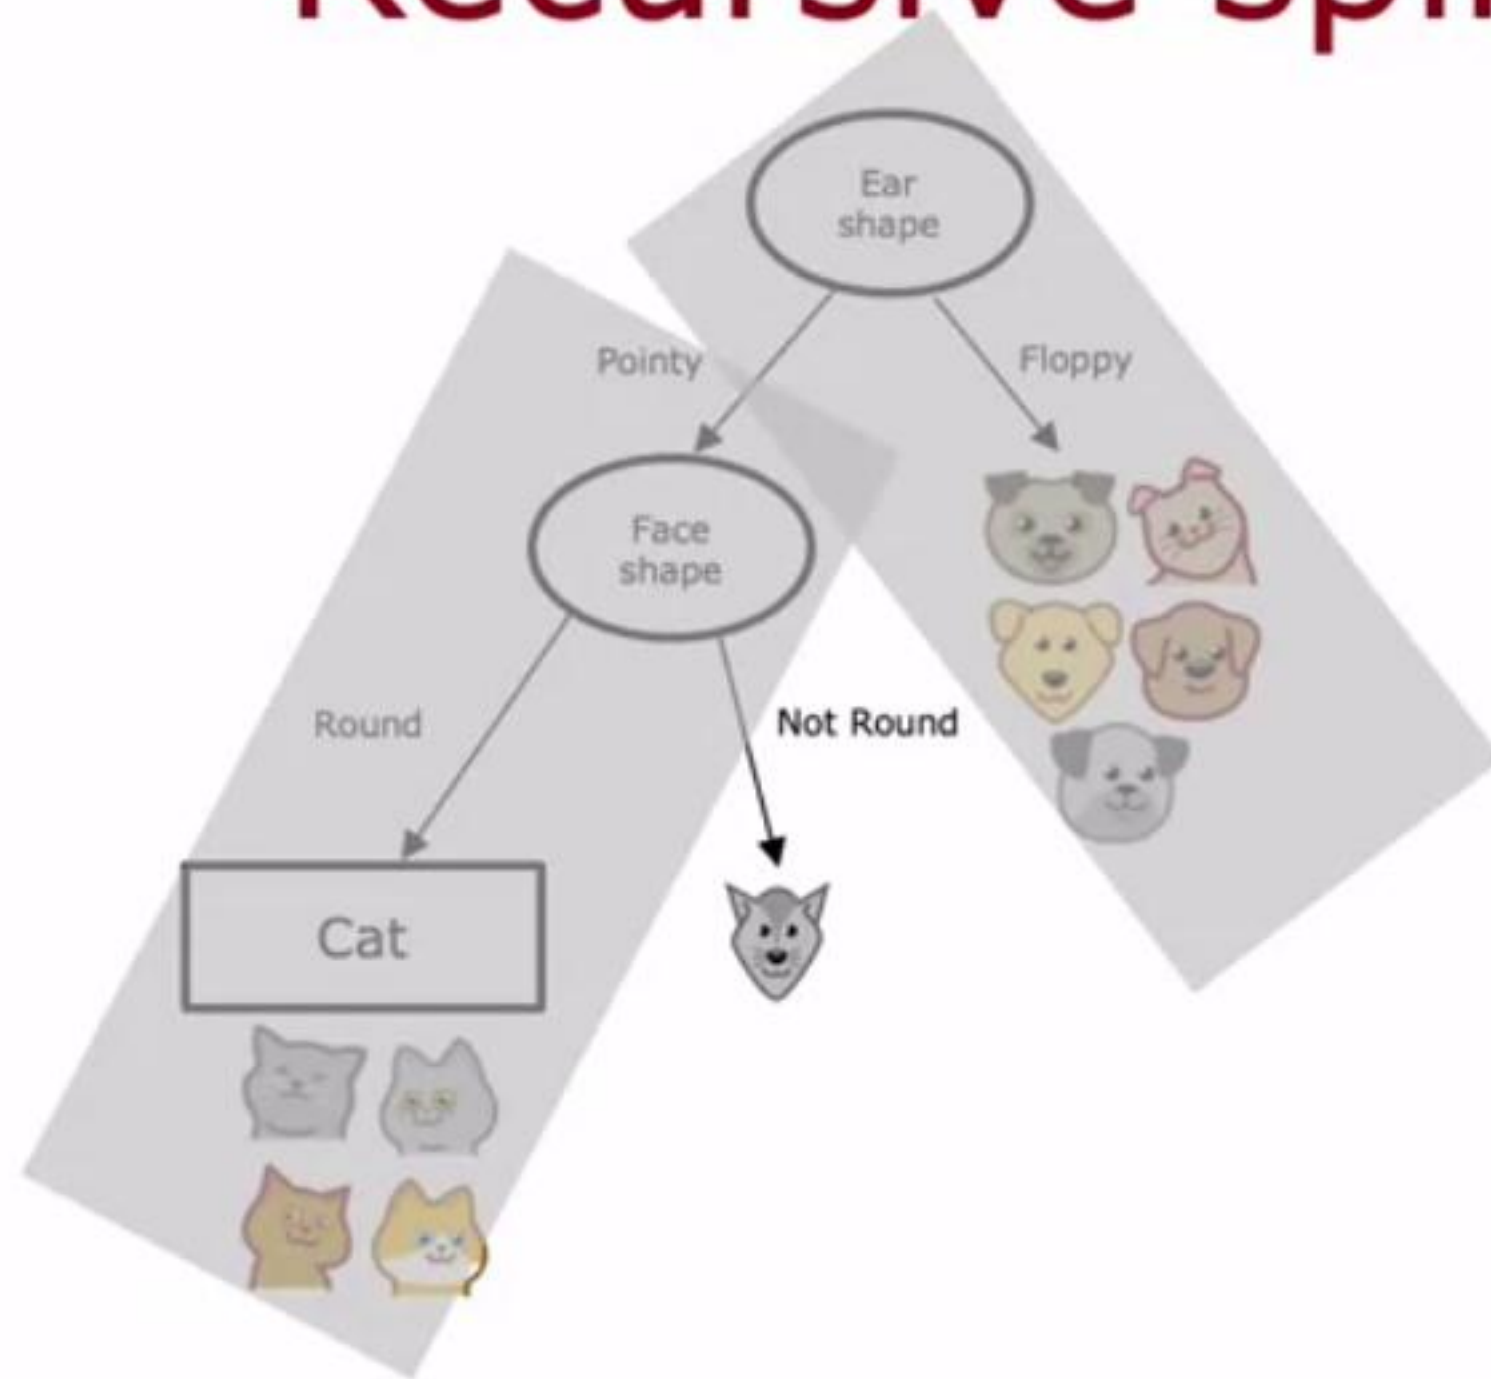




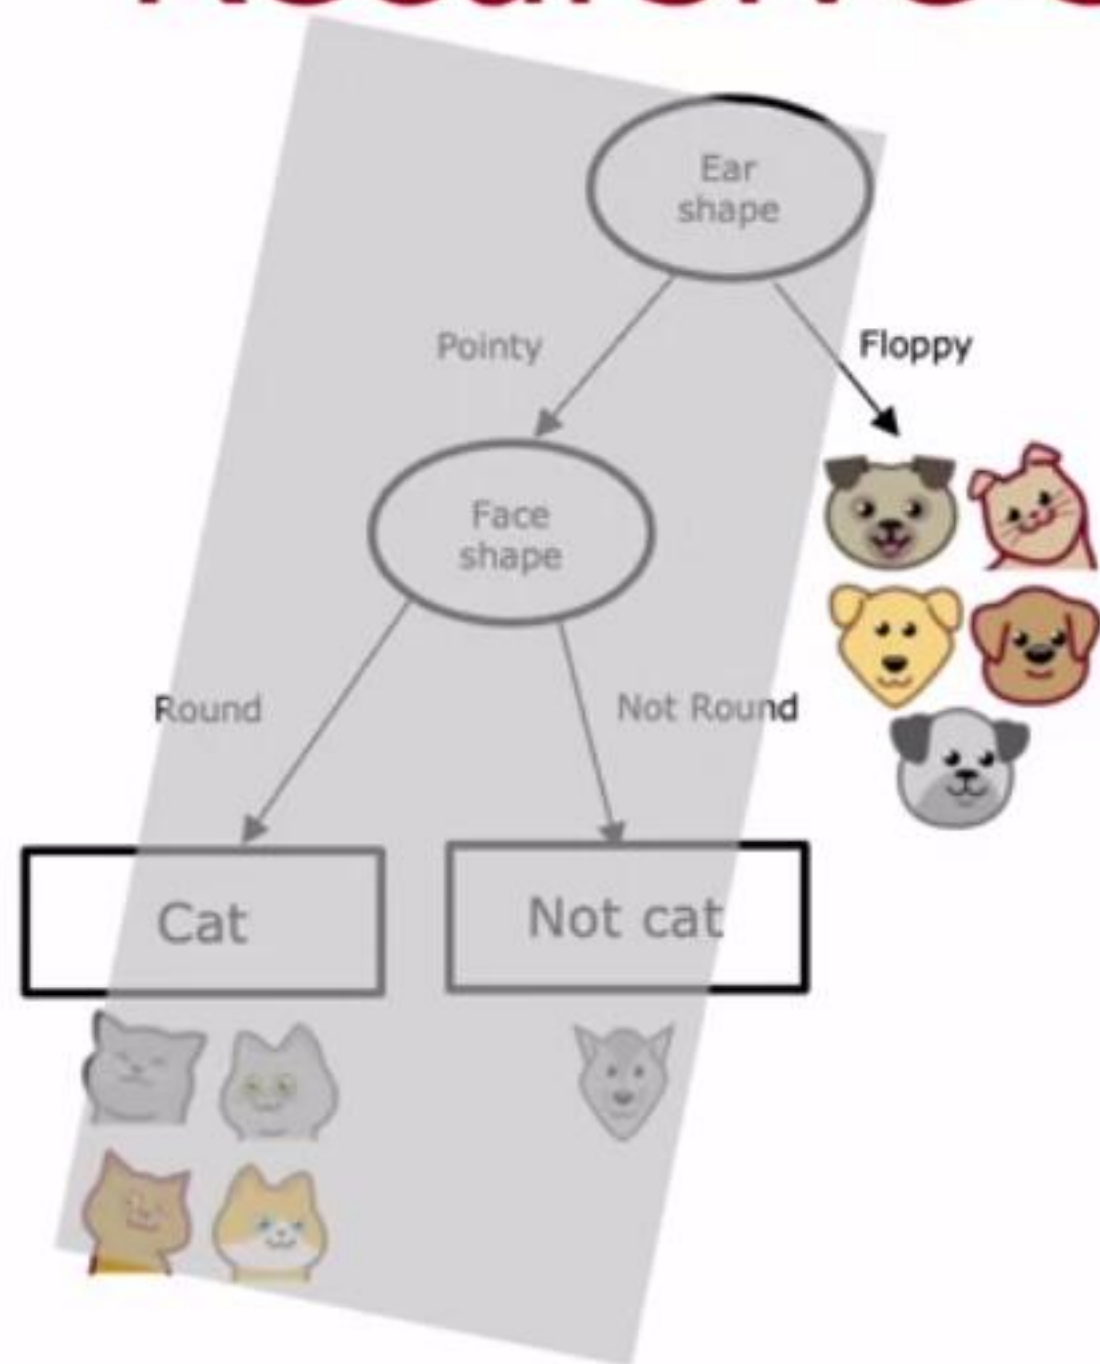
# Recursive splitting



# Recursive splitting

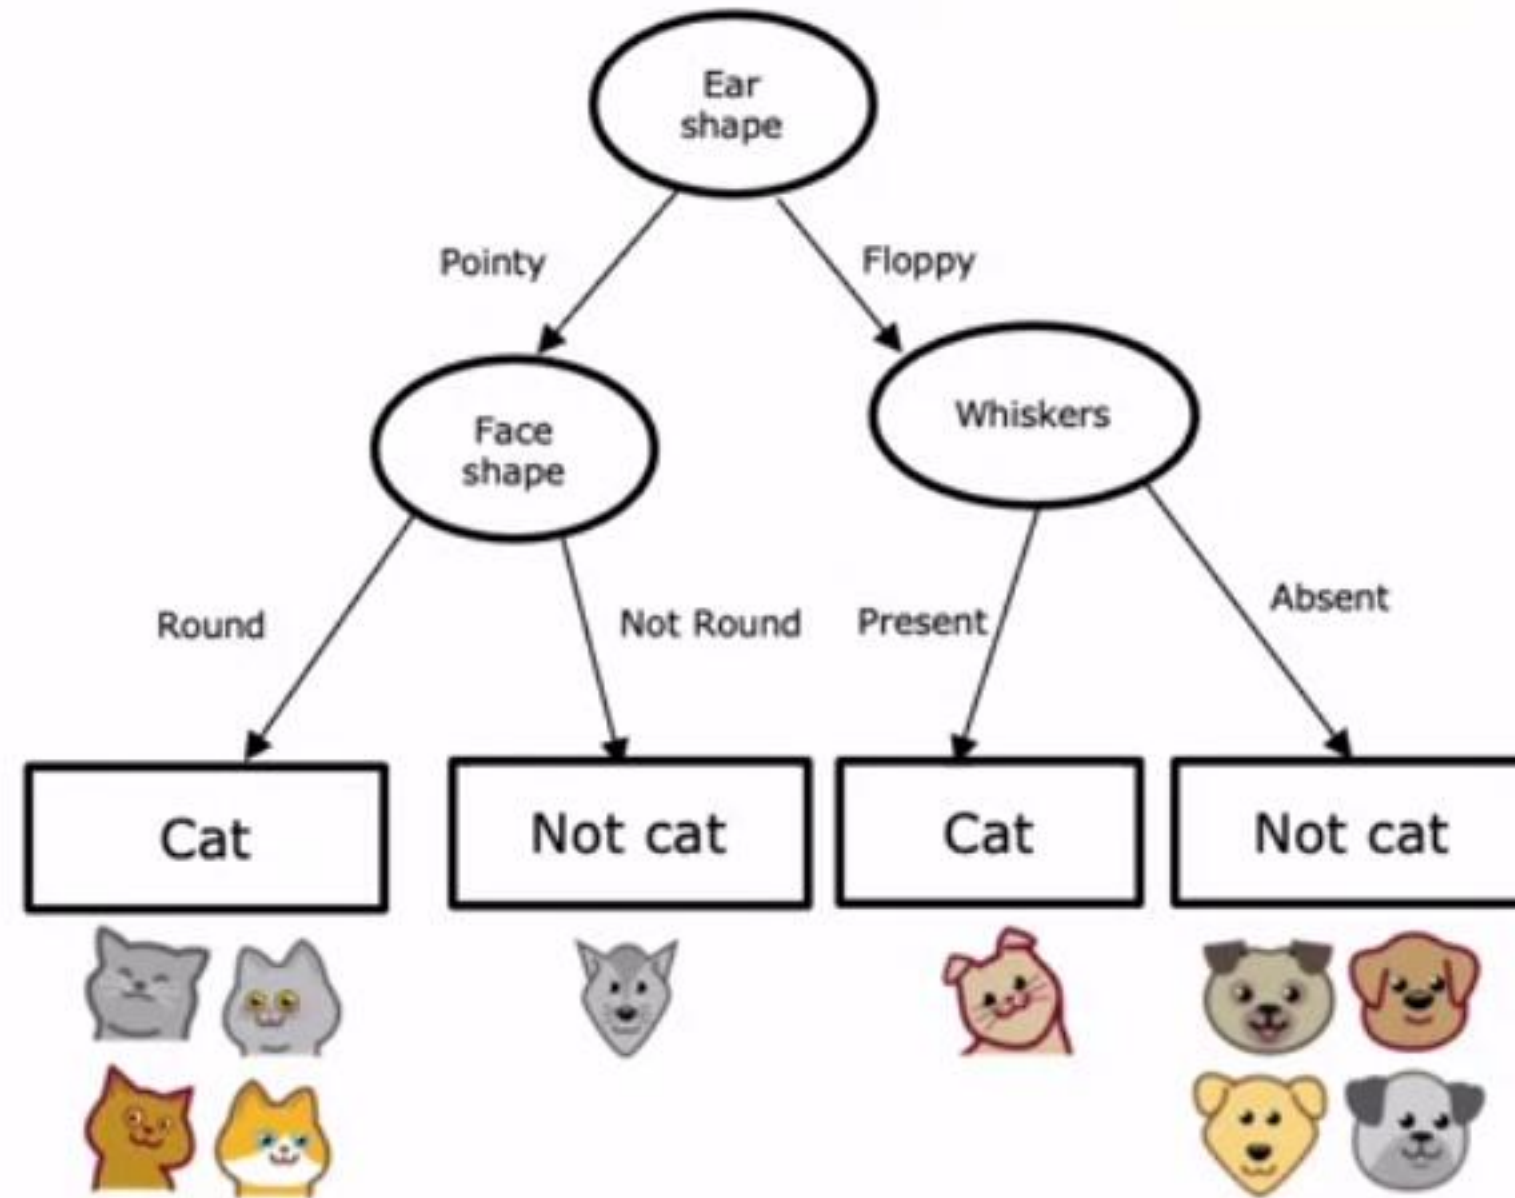


# Recursive splitting

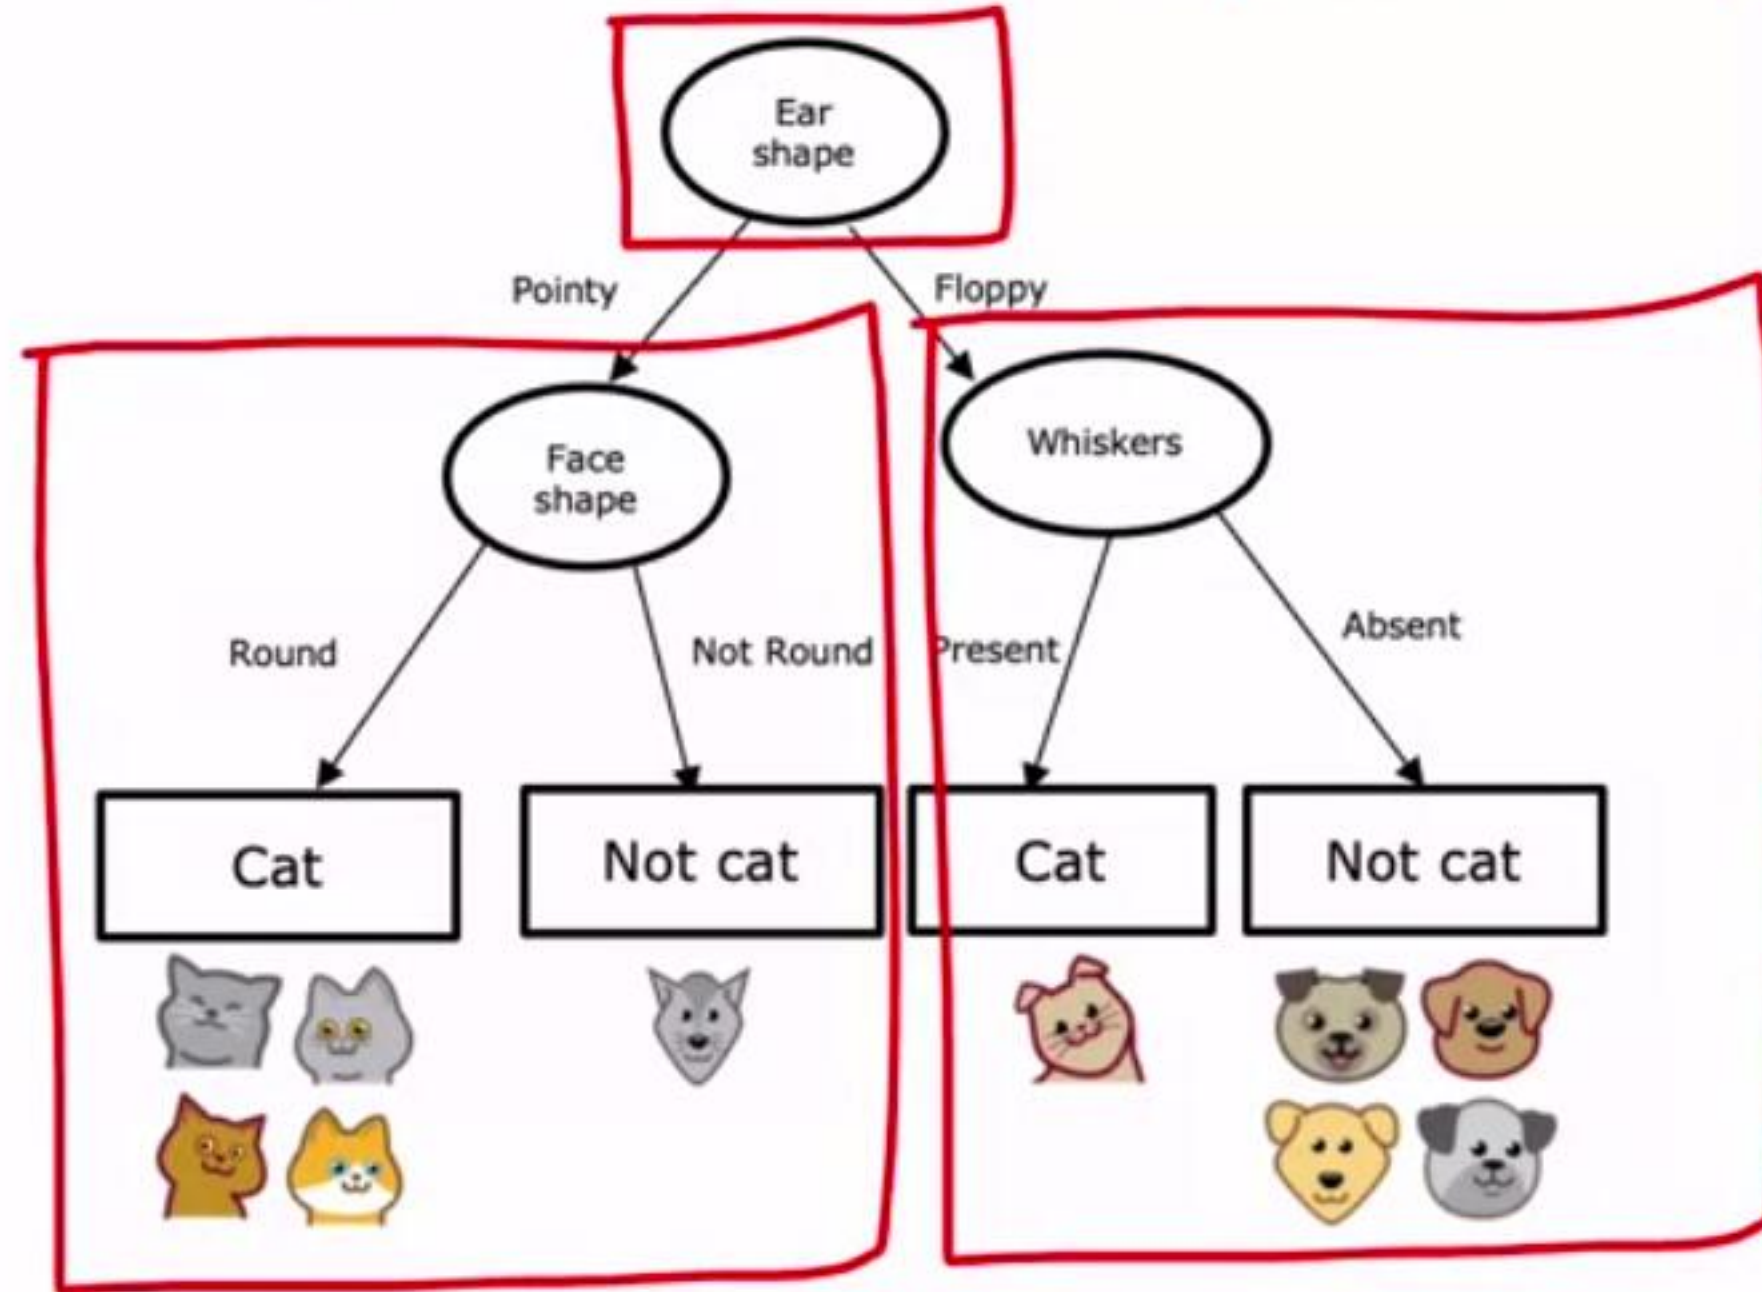




# Recursive splitting




# Recursive splitting



Recursive algorithm



Stanford  
ONLINE
















# Decision Tree Learning


Using one-hot encoding of  
categorical features







# Features with three possible values

	Ear shape ( $x_1$ )	Face shape ( $x_2$ )	Whiskers ( $x_3$ )	Cat ( $y$ )
	Pointy 	Round	Present	1
	Oval	Not round	Present	1
	Oval 	Round	Absent	0
	Pointy	Not round	Present	0
	Oval	Round	Present	1
	Pointy	Round	Absent	1
	Floppy 	Not round	Absent	0
	Oval	Round	Absent	1
	Floppy	Round	Absent	0
	Floppy	Round	Absent	0

3 possible values



# One hot encoding








Ear shape		Pointy ears	Floppy ears	Oval ears	Face shape	Whiskers	Cat
	Pointy	1	0	0	Round	Present	1
	Oval	0	0	1	Not round	Present	1
	Oval	0	0	1	Round	Absent	0
	Pointy	1	0	0	Not round	Present	0
	Oval	0	0	1	Round	Present	1
	Pointy	1	0	0	Round	Absent	1
	Floppy	0	1	0	Not round	Absent	0
	Oval	0	0	1	Round	Absent	1
	Floppy	0	1	0	Round	Absent	0
	Floppy	0	1	0	Round	Absent	0

# One hot encoding

If a categorical feature can take on  $k$  values, create  $k$  binary features (0 or 1 valued).




# One hot encoding and neural networks

	Pointy ears	Floppy ears	Round ears	Face shape	Whiskers	Cat
	1	0	0	<del>Round</del> 1	<del>Present</del> 1	1
	0	0	1	<del>Not round</del> 0	<del>Present</del> 1	1
	0	0	1	<del>Round</del> 1	<del>Absent</del> 0	0
	1	0	0	<del>Not round</del> 0	Present 1	0
	0	0	1	Round 1	Present 1	1
	1	0	0	Round 1	Absent 0	1
	0	1	0	<del>Not round</del> 0	Absent 0	1
	0	0	1	Round 1	Absent 0	1
	0	1	0	Round 1	Absent 0	1
	0	1	0	Round 1	Absent 0	1



Stanford  
ONLINE













# Decision Tree Learning

Continuous valued features




# Continuous features



	Ear shape	Face shape	Whiskers	Weight (lbs.)	Cat
	Pointy	Round	Present	7.2	1
	Floppy	Not round	Present	8.8	1
	Floppy	Round	Absent	15	0
	Pointy	Not round	Present	9.2	0
	Pointy	Round	Present	8.4	1
	Pointy	Round	Absent	7.6	1
	Floppy	Not round	Absent	11	0
	Pointy	Round	Absent	10.2	1
	Floppy	Round	Absent	18	0
	Floppy	Round	Absent	20	0




# Splitting on a continuous variable



$$H(0.5) - \left( \frac{2}{10} H\left(\frac{2}{2}\right) + \frac{8}{10} H\left(\frac{3}{8}\right) \right) = 0.24$$


$$H(0.5) - \left( \frac{4}{10} H\left(\frac{4}{4}\right) + \frac{6}{10} H\left(\frac{1}{6}\right) \right) = 0.61$$

$$H(0.5) - \left( \frac{7}{10} H\left(\frac{5}{7}\right) + \frac{3}{10} H\left(\frac{0}{3}\right) \right) = 0.40$$















Stanford  
ONLINE





# Decision Tree Learning

Regression Trees (optional)

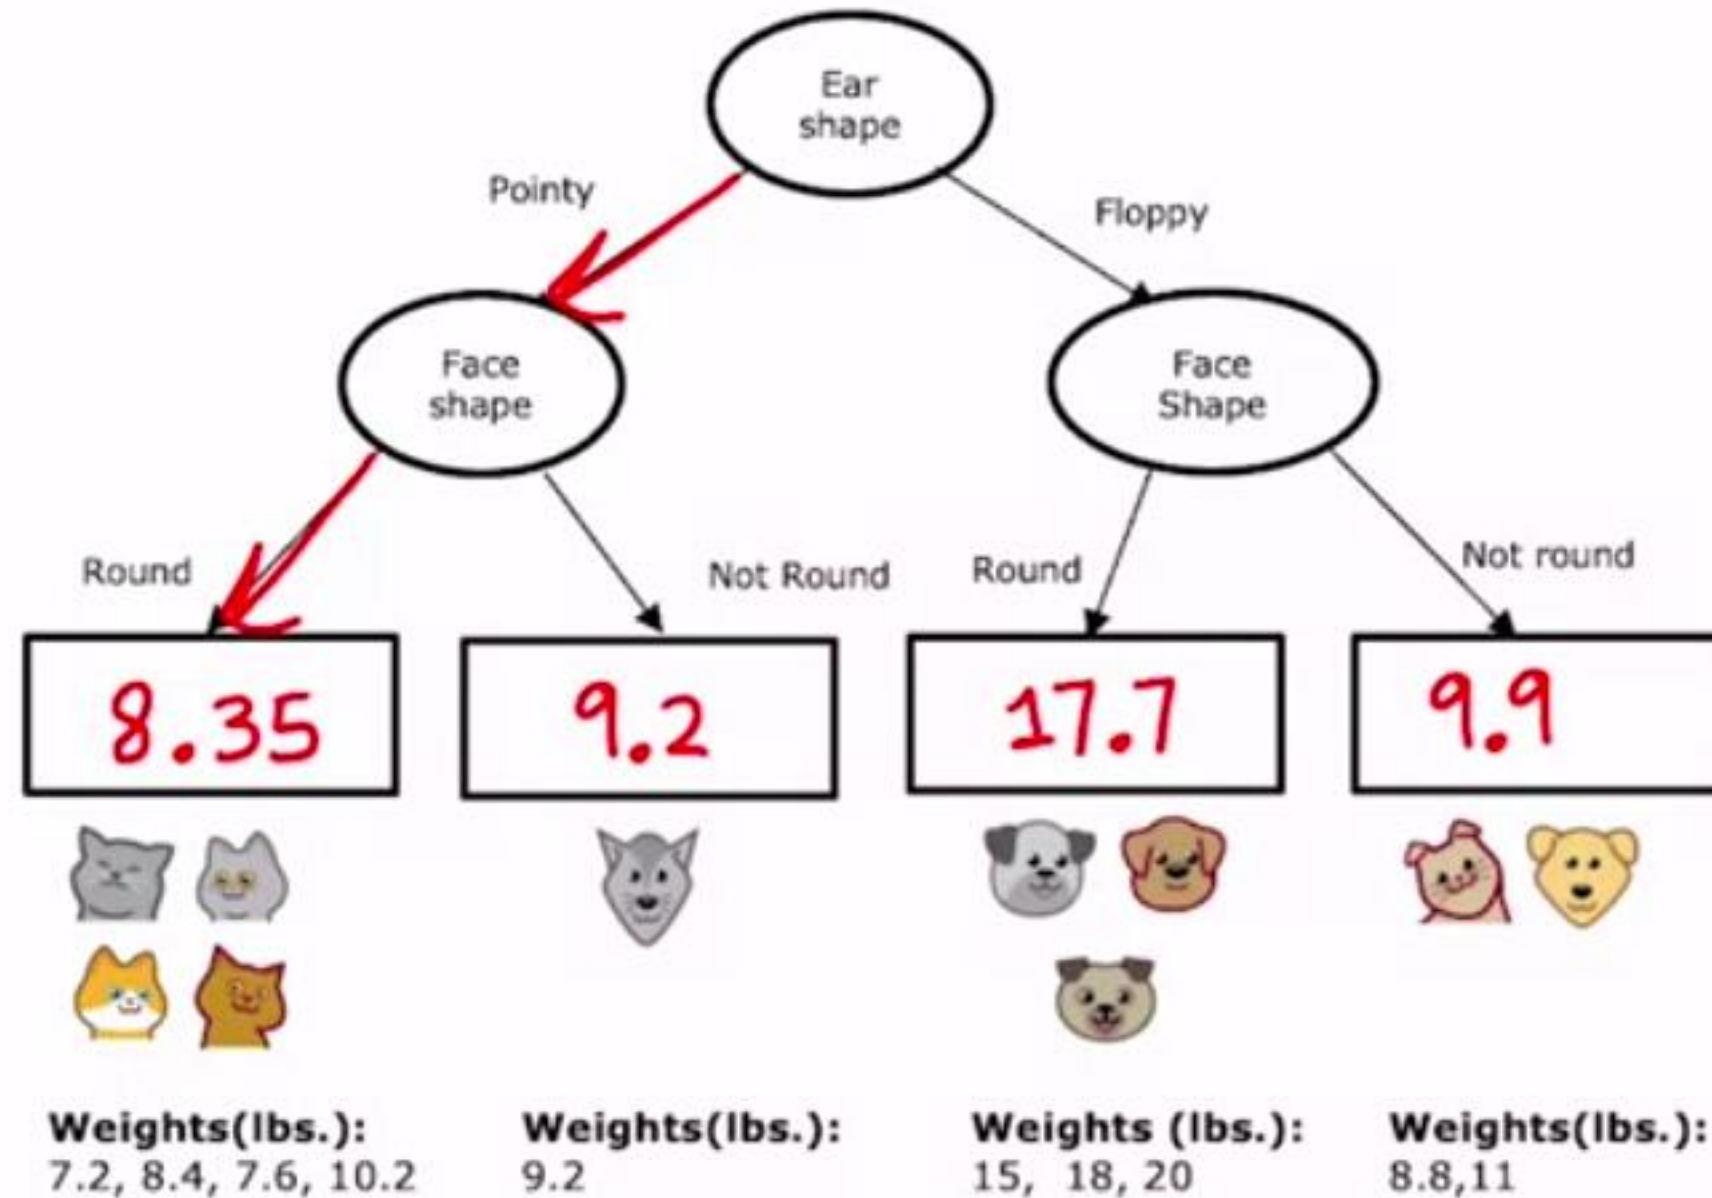
# Regression with Decision Trees: Predicting a number

	Ear shape	Face shape	Whiskers	Weight (lbs.)
	Pointy	Round	Present	7.2
	Floppy	Not round	Present	8.8
	Floppy	Round	Absent	15
	Pointy	Not round	Present	9.2
	Pointy	Round	Present	8.4
	Pointy	Round	Absent	7.6
	Floppy	Not round	Absent	11
	Pointy	Round	Absent	10.2
	Floppy	Round	Absent	18
	Floppy	Round	Absent	20

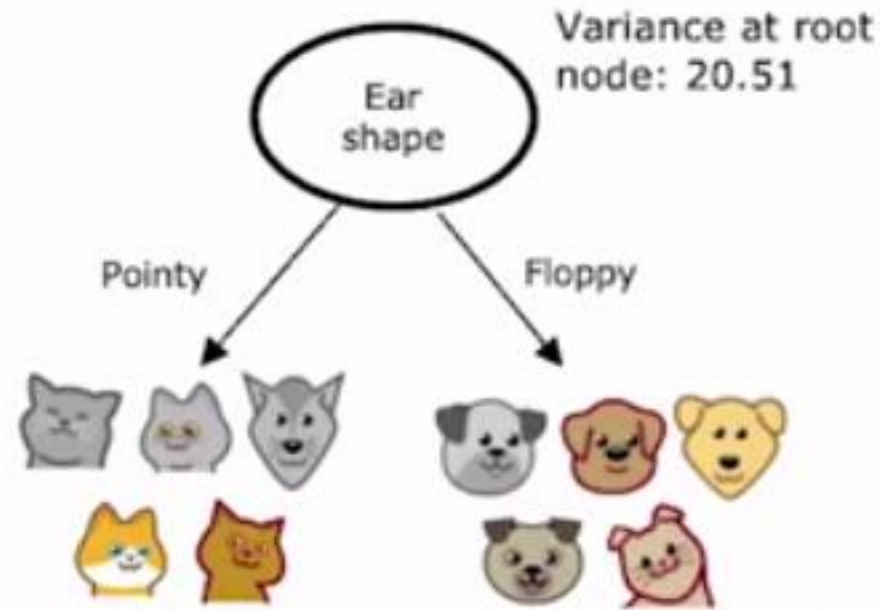
 



# Regression with Decision Trees



# Choosing a split



Weights: 7.2, 9.2, 8.4, 7.6, 10.2

Weights: 8.8, 15, 11, 18, 20

Variance: 1.47


Variance: 21.87

$$w^{\text{left}} = 5/10$$

$$w^{\text{right}} = 5/10$$

$$20.51 - \left( \frac{5}{10} * 1.47 + \frac{5}{10} * 21.87 \right)$$

$$= 8.84 \leftarrow$$



Weights: 7.2, 15, 8.4, 7.6, 10.2, 18, 20

Weights: 8.8, 9.2, 11

Variance: 27.80

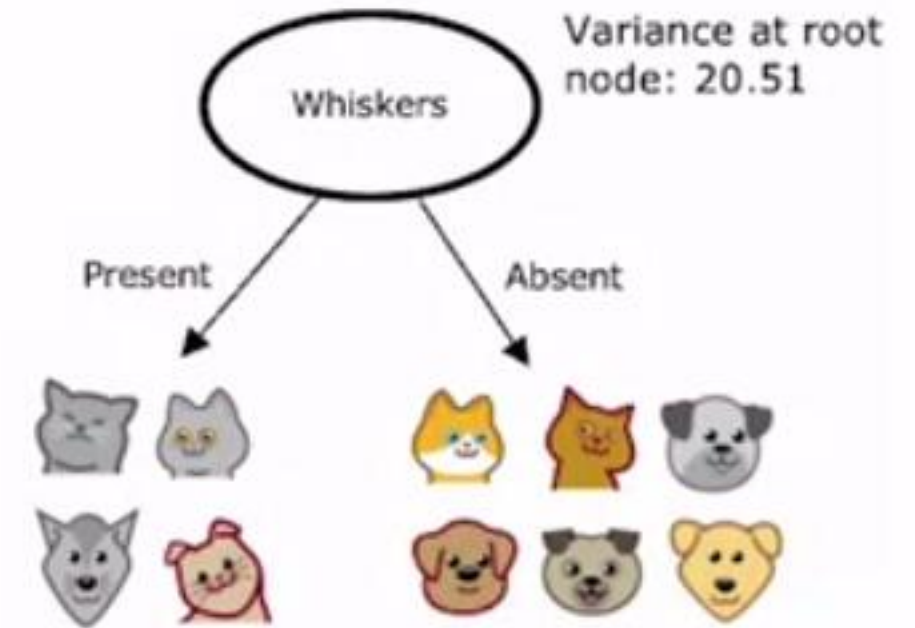
Variance: 1.37

$$w^{\text{left}} = 7/10$$

$$w^{\text{right}} = 3/10$$

$$20.51 - \left( \frac{7}{10} * 27.80 + \frac{3}{10} * 1.37 \right)$$

$$= 0.64$$



Weights: 7.2, 8.8, 9.2, 8.4

Weights: 15, 7.6, 11, 10.2, 18, 20

Variance: 0.75

Variance: 23.32

$$w^{\text{left}} = 4/10$$

$$w^{\text{right}} = 6/10$$

$$20.51 - \left( \frac{4}{10} * 0.75 + \frac{6}{10} * 23.32 \right)$$

$$= 6.22$$





# Tree ensembles

Using multiple decision trees



# Trees are highly sensitive to small changes of the data

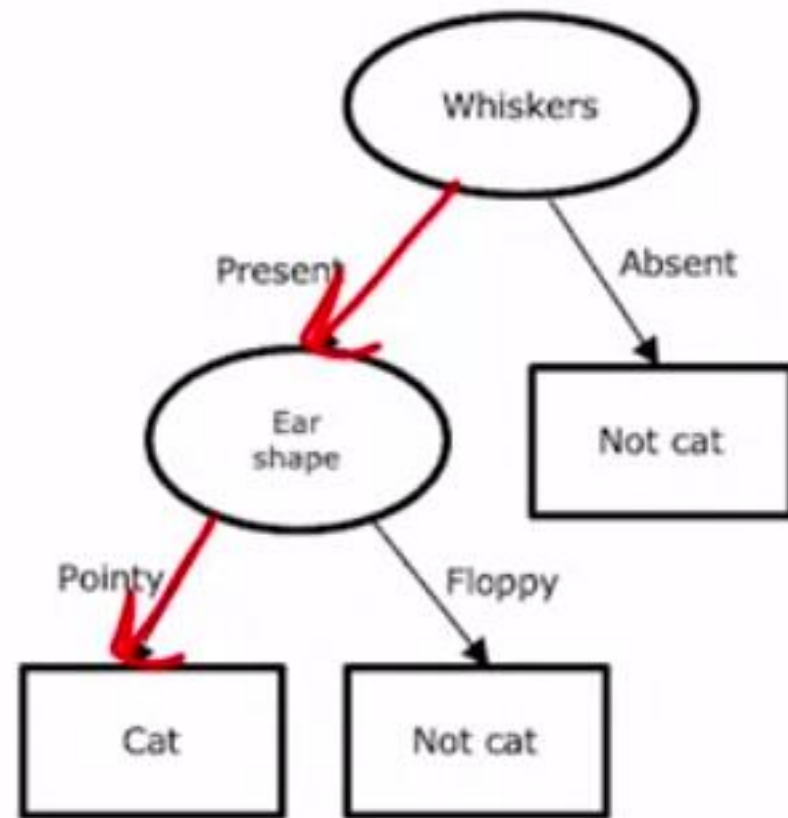


# Tree ensemble

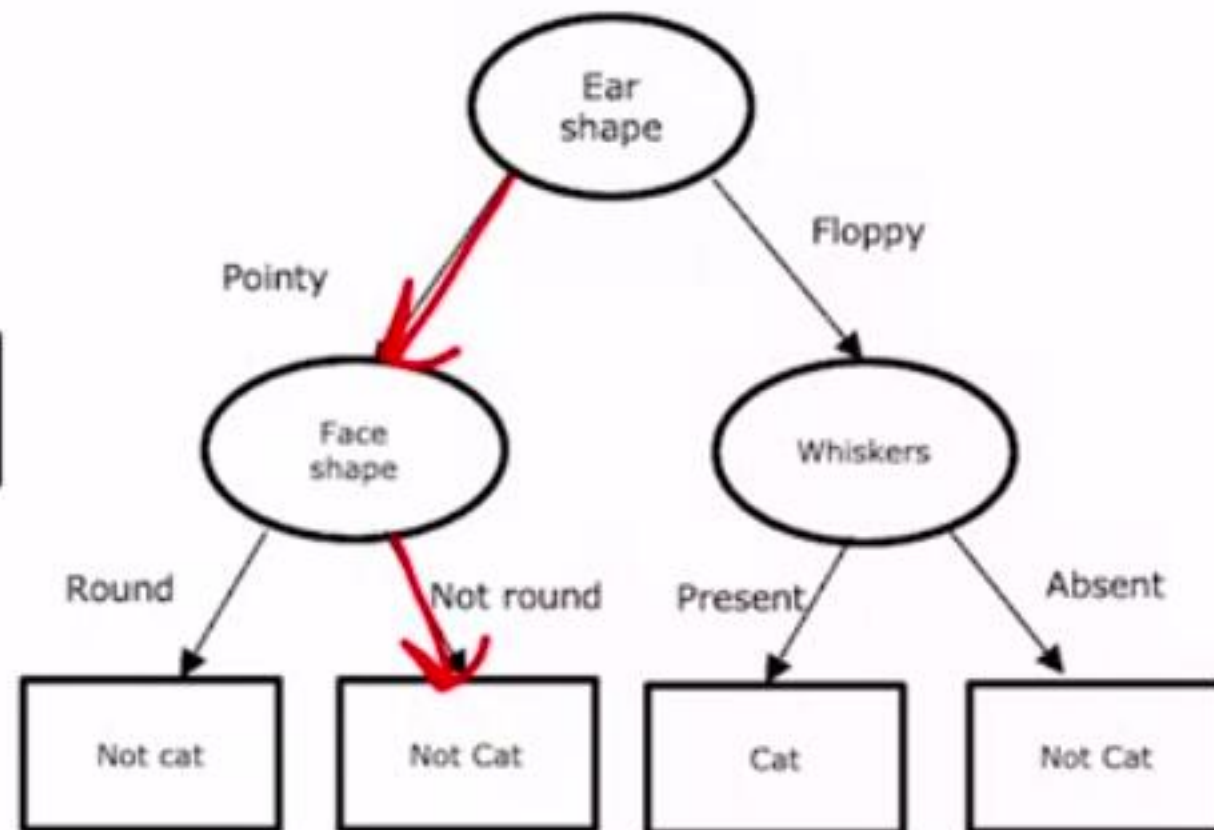
New test example



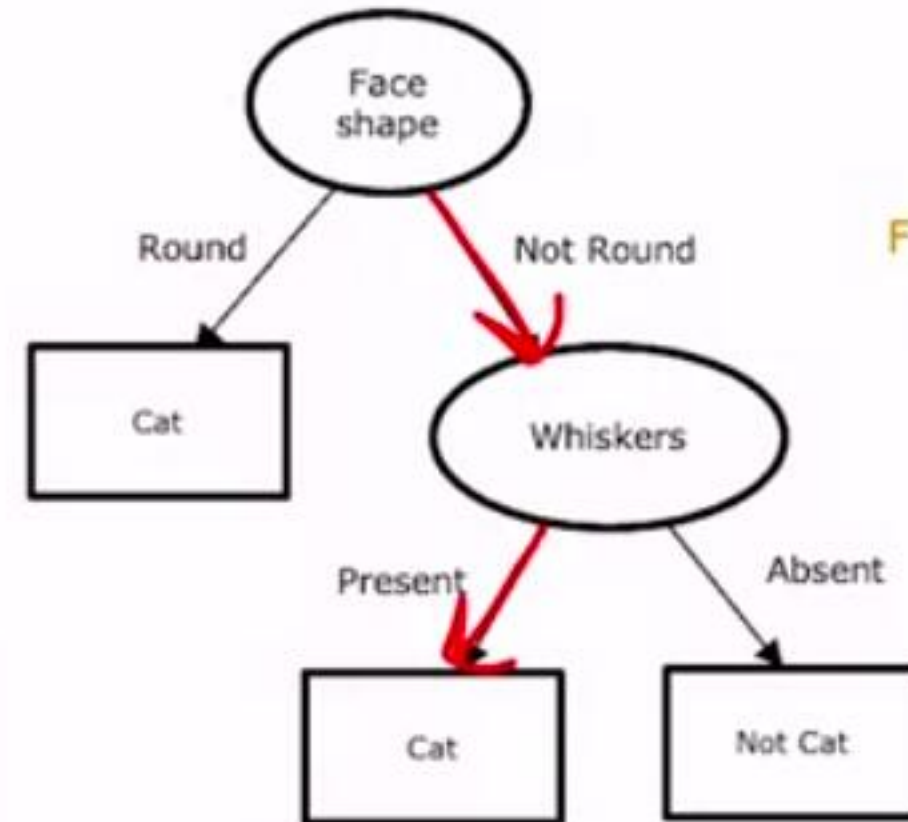
Ear shape: Pointy  
Face shape: Not Round  
Whiskers: Present



Prediction: Cat



Prediction: Not cat



Prediction: Cat





# Tree ensembles


Sampling with replacement



# Sampling with replacement











Tokens    

Sampling with replacement:



# Sampling with replacement



	Ear shape	Face shape	Whiskers	Cat
	Pointy	Round	Present	1
	Floppy	Not round	Absent	0
	Pointy	Round	Absent	1
	Pointy	Not round	Present	0
	Floppy	Not round	Absent	0
	Pointy	Round	Absent	1
	Pointy	Round	Present	1
	Floppy	Not round	Present	1
	Floppy	Round	Absent	0
	Pointy	Round	Absent	1





Stanford  
ONLINE



# Tree ensembles

## Random forest algorithm



# Generating a tree sample

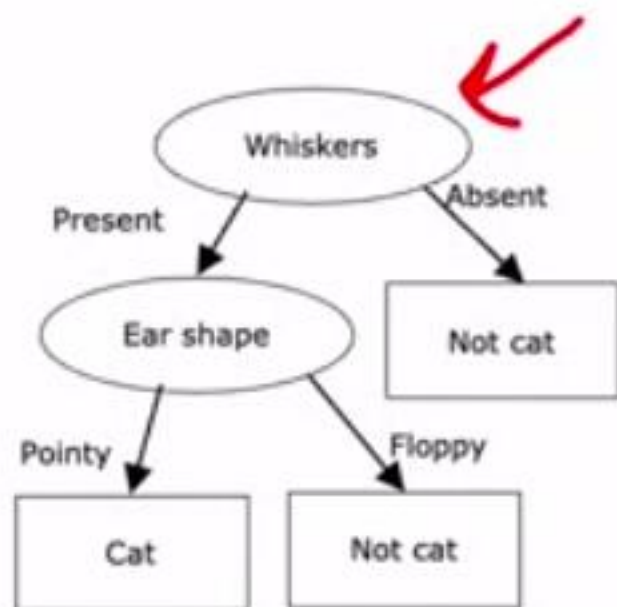
Given training set of size  $m$

For  $b = 1$  to  $B$ :


Use sampling with replacement to create a new training set of size  $m$

Train a decision tree on the new dataset

Ear shape	Face shape	Whiskers	Cat
Pointy	Round	Present	Yes
Floppy	Round	Absent	No
Floppy	Round	Absent	No
Pointy	Round	Present	Yes
Pointy	Not Round	Present	Yes
Floppy	Round	Absent	No
Floppy	Round	Present	Yes
Floppy	Round	Present	Yes
Pointy	Not Round	Absent	No
Pointy	Not Round	Absent	No
Pointy	Not Round	Present	Yes



Ear shape	Face shape	Whiskers	Cat
Pointy	Round	Present	Yes
Pointy	Round	Absent	Yes
Floppy	Not Round	Absent	No
Floppy	Not Round	Absent	No
Pointy	Round	Absent	Yes
Floppy	Round	Absent	No
Floppy	Round	Absent	No
Floppy	Round	Absent	No
Floppy	Round	Absent	No
Floppy	Not Round	Absent	No
Pointy	Round	Absent	No
Pointy	Round	Present	Yes



...

Bagged decision tree

# Randomizing the feature choice

At each node, when choosing a feature to use to split, if  $n$  features are available, pick a random subset of  $k < n$  features and allow the algorithm to only choose from that subset of features.

$$k = \sqrt{n}$$

Random forest algorithm





Stanford  
ONLINE



# Tree ensembles

# XGBoost



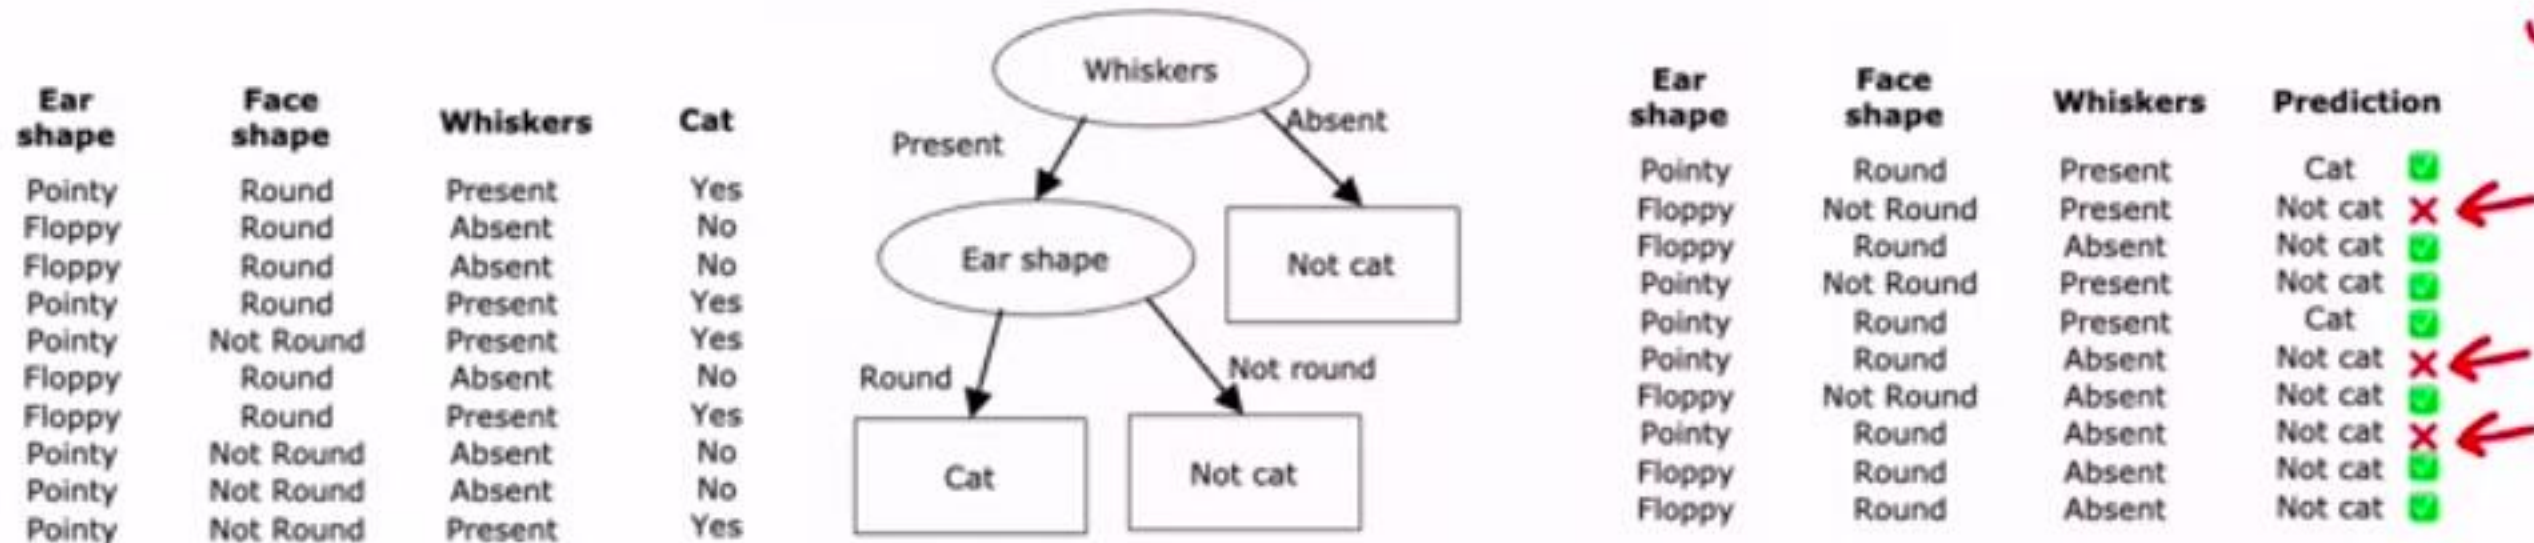
# Boosted trees intuition

Given training set of size  $m$

For  $b = 1$  to  $B$ :

Use sampling with replacement to create a new training set of size  $m$   
But instead of picking from all examples with equal ( $1/m$ ) probability, make it more likely to pick misclassified examples from previously trained trees

Train a decision tree on the new dataset



# XGBoost (eXtreme Gradient Boosting)

- Open source implementation of boosted trees
- Fast efficient implementation
- Good choice of default splitting criteria and criteria for when to stop splitting
- Built in regularization to prevent overfitting
- Highly competitive algorithm for machine learning competitions (eg: Kaggle competitions)



# Using XGBoost

## Classification

```
→ from xgboost import XGBClassifier  
  
→ model = XGBClassifier()  
  
→ model.fit(X_train, y_train)  
→ y_pred = model.predict(X_test)
```

## Regression

```
from xgboost import XGBRegressor  
  
model = XGBRegressor()  
  
model.fit(X_train, y_train)  
y_pred = model.predict(X_test)
```





# Conclusion

**When to use decision trees**

# Decision Trees vs Neural Networks

## Decision Trees and Tree ensembles

- Works well on tabular (structured) data
- Not recommended for unstructured data (images, audio, text)
- Fast
- Small decision trees may be human interpretable

## Neural Networks

- Works well on all types of data, including tabular (structured) and unstructured data
- May be slower than a decision tree
- Works with transfer learning
- When building a system of multiple models working together, it might be easier to string together multiple neural networks