

Stanford  
ONLINE

DeepLearning.AI



# Linear Regression with Multiple Variables

---

## Multiple Features

# Multiple features (variables)

one  
feature



| Size in feet <sup>2</sup> ( $x$ ) | Price (\$) in 1000's ( $y$ ) |
|-----------------------------------|------------------------------|
| 2104                              | 400                          |
| 1416                              | 232                          |
| 1534                              | 315                          |
| 852                               | 178                          |
| ...                               | ...                          |



$$f_{w,b}(x) = wx + b$$



# Multiple features (variables)

|       | Size in feet <sup>2</sup><br>$x_1$ | Number of bedrooms<br>$x_2$ | Number of floors<br>$x_3$ | Age of home in years<br>$x_4$ | Price (\$) in \$1000's |
|-------|------------------------------------|-----------------------------|---------------------------|-------------------------------|------------------------|
|       | 2104                               | 5                           | 1                         | 45                            | 460                    |
| $i=2$ | 1416                               | 3                           | 2                         | 40                            | 232                    |
|       | 1534                               | 3                           | 2                         | 30                            | 315                    |
|       | 852                                | 2                           | 1                         | 36                            | 178                    |
|       | ...                                | ...                         | ...                       | ...                           | ...                    |

$$j = 1 \dots 4$$

$$n = 4$$

$x_j = j^{\text{th}}$  feature

$n$  = number of features

$\vec{x}^{(i)}$  = features of  $i^{\text{th}}$  training example

$x_j^{(i)}$  = value of feature  $j$  in  $i^{\text{th}}$  training example

$$\vec{x}^{(2)} = [1416 \ 3 \ 2 \ 40]$$

$$x_3^{(2)} = 2$$

# Model:

Previously:  $f_{w,b}(x) = wx + b$

$$f_{w,b}(x) = w_1 x_1 + w_2 x_2 + w_3 x_3 + w_4 x_4 + b$$

example

$$f_{w,b}(x) = 0.1 \underset{\substack{\uparrow \\ \text{size}}}{x_1} + 4 \underset{\substack{\uparrow \\ \text{\# bedrooms}}}{x_2} + 10 \underset{\substack{\uparrow \\ \text{\# floors}}}{x_3} + -2 \underset{\substack{\uparrow \\ \text{years}}}{x_4} + 80 \underset{\substack{\uparrow \\ \text{base price}}}{}$$



$$f_{\vec{w},b}(\vec{x}) = w_1 x_1 + w_2 x_2 + \dots + w_n x_n + b$$

$$\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n] \quad \text{parameters of the model}$$

$b$  is a number

vector  $\vec{x} = [x_1 \ x_2 \ x_3 \ \dots \ x_n]$

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1 x_1 + w_2 x_2 + w_3 x_3 + \dots + w_n x_n + b$$

dot product

multiple linear regression

(not multivariate regression)

Stanford  
ONLINE

DeepLearning.AI



# Linear Regression with Multiple Variables

---

## Vectorization Part 1



Stanford  
ONLINE

DeepLearning.AI



# Linear Regression with Multiple Variables

---

## Vectorization Part 1



## Parameters and features

$$\vec{w} = [w_1 \ w_2 \ w_3] \quad n=3$$

$b$  is a number

$$\vec{x} = [x_1 \ x_2 \ x_3]$$

linear algebra: count from 1

$w[0]$   $w[1]$   $w[2]$

NumPy 

```
w = np.array([1.0, 2.5, -3.3])
```

```
b = 4
```

$x[0]$   $x[1]$   $x[2]$

```
x = np.array([10, 20, 30])
```

code: count from 0

## Without vectorization $n=100,000$

$$f_{\vec{w},b}(\vec{x}) = w_1x_1 + w_2x_2 + w_3x_3 + b$$

```
f = w[0] * x[0] +  
     w[1] * x[1] +  
     w[2] * x[2] + b
```



## Without vectorization

$$f_{\vec{w},b}(\vec{x}) = \left( \sum_{j=1}^n w_j x_j \right) + b$$

$\sum_{j=1}^n \rightarrow j=1 \dots n$   
 $1, 2, 3$

$\text{range}(0, n) \rightarrow j=0 \dots n-1$

```
f = 0  
for j in range(0, n):  
    f = f + w[j] * x[j]  
f = f + b
```



## Vectorization

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$$

```
f = np.dot(w, x) + b
```





Stanford  
ONLINE

DeepLearning.AI



# Linear Regression with Multiple Variables

---

## Vectorization Part 2

## Without vectorization

```
for j in range(0,16):  
    f = f + w[j] * x[j]
```

$t_0$

$$f + w[0] * x[0]$$

$t_1$

$$f + w[1] * x[1]$$

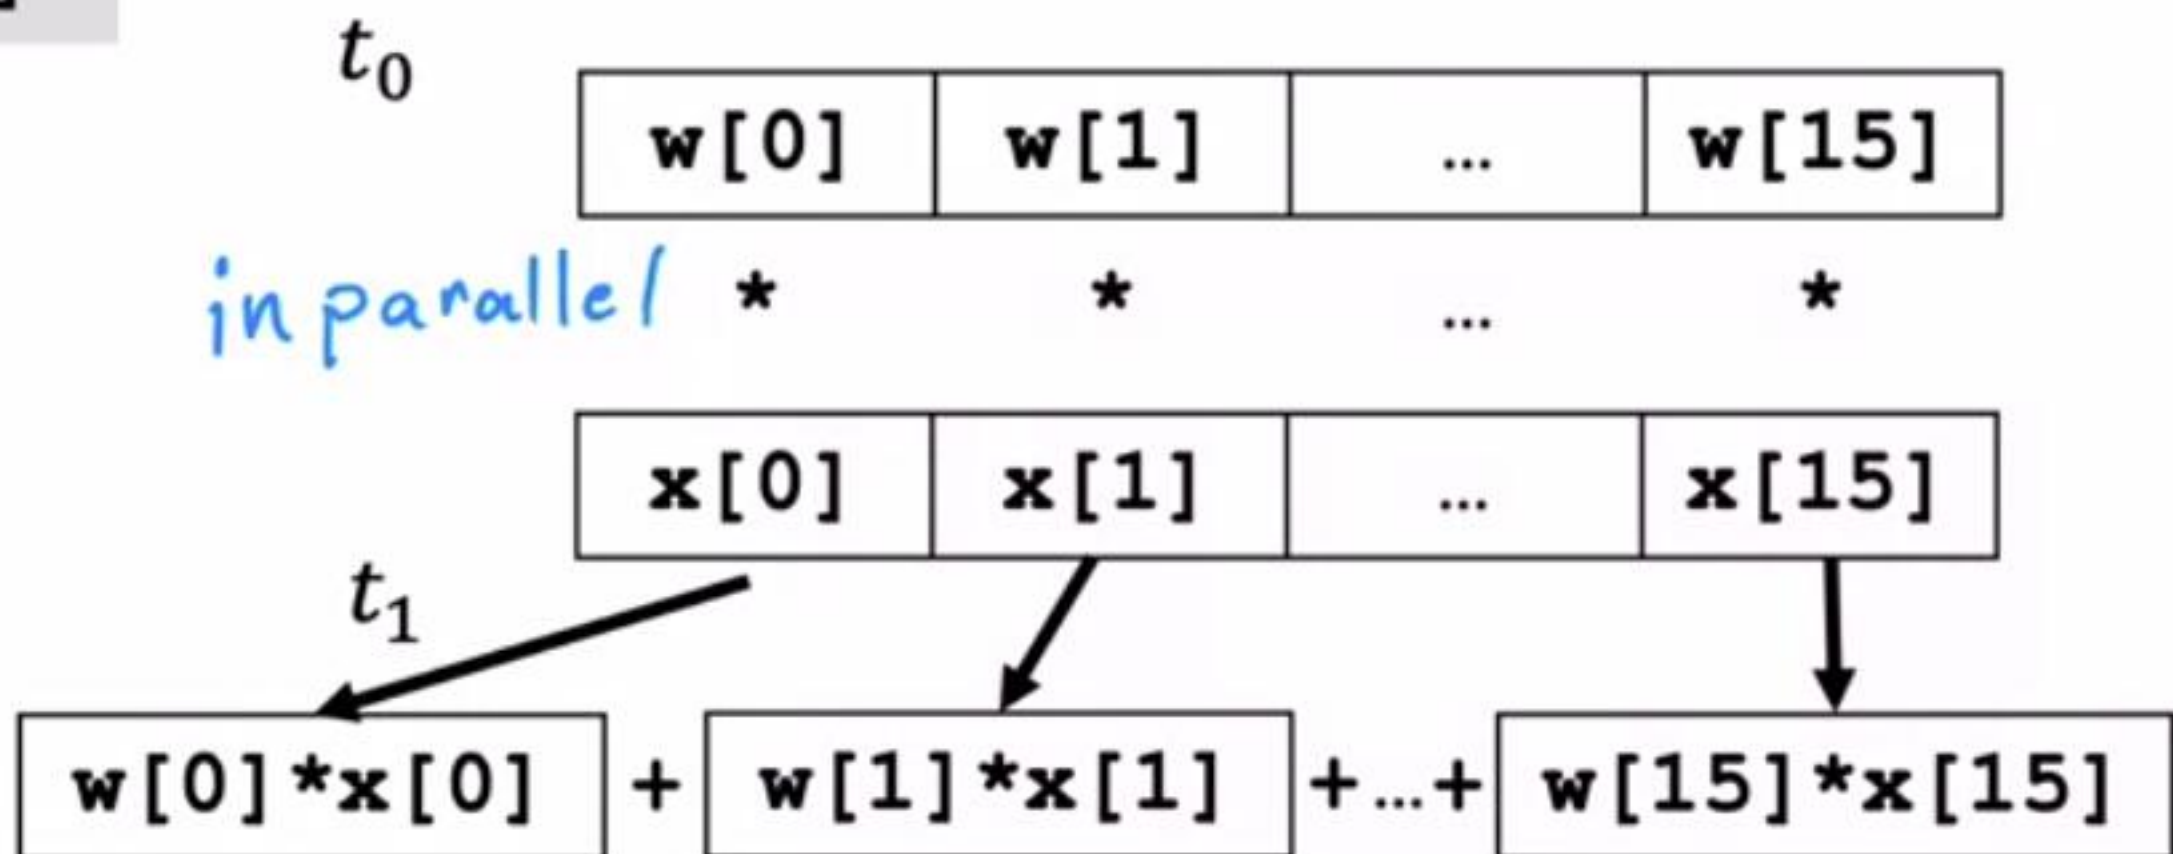
...

$t_{15}$

$$f + w[15] * x[15]$$

## Vectorization

```
np.dot(w,x)
```



*efficient → scale to large datasets*



## Gradient descent

$$\vec{w} = (w_1 \quad w_2 \quad \dots \quad w_{16})$$

~~b~~ parameters

derivatives  $\vec{d} = (d_1 \quad d_2 \quad \dots \quad d_{16})$

```
w = np.array([0.5, 1.3, ... 3.4])
```

```
d = np.array([0.3, 0.2, ... 0.4])
```

compute  $w_j = w_j - \underbrace{0.1}_{\text{learning rate } \alpha} d_j$  for  $j = 1 \dots 16$

### Without vectorization

$$w_1 = w_1 - 0.1d_1$$

$$w_2 = w_2 - 0.1d_2$$

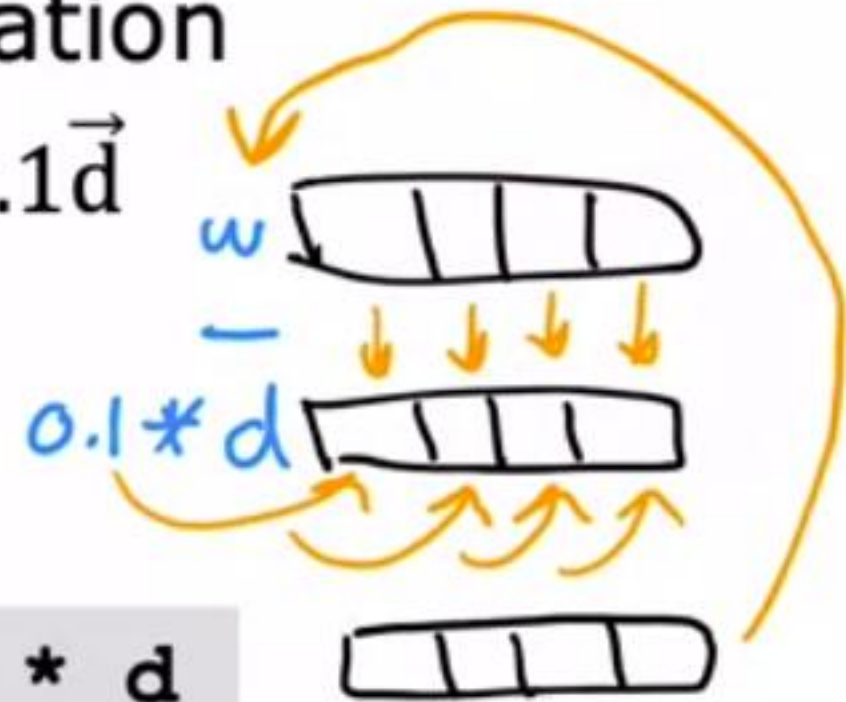
$$\vdots$$

$$w_{16} = w_{16} - 0.1d_{16}$$

```
for j in range(0,16):  
    w[j] = w[j] - 0.1 * d[j]
```

### With vectorization

$$\vec{w} = \vec{w} - 0.1\vec{d}$$



```
w = w - 0.1 * d
```

Stanford  
ONLINE

DeepLearning.AI



# Linear Regression with Multiple Variables

---

## Gradient Descent for Multiple Regression



## Previous notation

Parameters

$$w_1, \dots, w_n$$
$$b$$

Model

$$f_{\vec{w},b}(\vec{x}) = w_1 x_1 + \dots + w_n x_n + b$$

Cost function

$$J(\underbrace{w_1, \dots, w_n}_{\text{vector}}, b)$$

Gradient descent

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(\underbrace{w_1, \dots, w_n}_{\text{vector}}, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(\underbrace{w_1, \dots, w_n}_{\text{vector}}, b)$$

}

## Vector notation

↖ vector of length n

$$\vec{w} = [w_1 \quad \dots \quad w_n]$$

b still a number

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$$

↗ dot product

$$J(\underbrace{\vec{w}}_{\text{vector}}, b)$$

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(\underbrace{\vec{w}}_{\text{vector}}, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(\underbrace{\vec{w}}_{\text{vector}}, b)$$

}

# Gradient descent

One feature

repeat {

$$\underline{w} = w - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) \underline{x^{(i)}}$$

$$\searrow \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$

simultaneously update  $w, b$

}

$n$  features ( $n \geq 2$ )

repeat {

$$j=1 \quad \underline{w_1} = w_1 - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\bar{w},b}(\bar{x}^{(i)}) - y^{(i)}) \underline{x_1^{(i)}}$$

$\vdots$

$j=n$

$$w_n = w_n - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\bar{w},b}(\bar{x}^{(i)}) - y^{(i)}) x_n^{(i)}$$

$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\bar{w},b}(\bar{x}^{(i)}) - y^{(i)})$$

simultaneously update

$w_j$  (for  $j = 1, \dots, n$ ) and  $b$

}



# An alternative to gradient descent

## → Normal equation

- Only for linear regression
- Solve for  $w$ ,  $b$  without iterations

### Disadvantages

- Doesn't generalize to other learning algorithms.
- Slow when number of features is large ( $> 10,000$ )

### What you need to know

- Normal equation method may be used in machine learning libraries that implement linear regression.
- Gradient descent is the recommended method for finding parameters  $w, b$

Stanford  
ONLINE

DeepLearning.AI



# Practical Tips for Linear Regression

---

## Feature Scaling Part 1



# Feature and parameter values

$$\widehat{price} = w_1 x_1 + w_2 x_2 + b$$

$x_1$ : size (feet<sup>2</sup>) range: 300 – 2,000 large  
 $x_2$ : # bedrooms range: 0 – 5 small

size    # bedrooms

House:  $x_1 = 2000$ ,  $x_2 = 5$ ,  $price = \$500k$     one training example

size of the parameters  $w_1, w_2$ ?

$w_1 = 50$ ,  $w_2 = 0.1$ ,  $b = 50$

$w_1 = 0.1$ ,  $w_2 = 50$ ,  $b = 50$   
small    large

$$\widehat{price} = \underbrace{50 * 2000}_{100,000k} + \underbrace{0.1 * 5}_{0.5k} + \underbrace{50}_{50k}$$

$$\widehat{price} = \underbrace{0.1 * 2000k}_{200k} + \underbrace{50 * 5}_{250k} + \underbrace{50}_{50k}$$

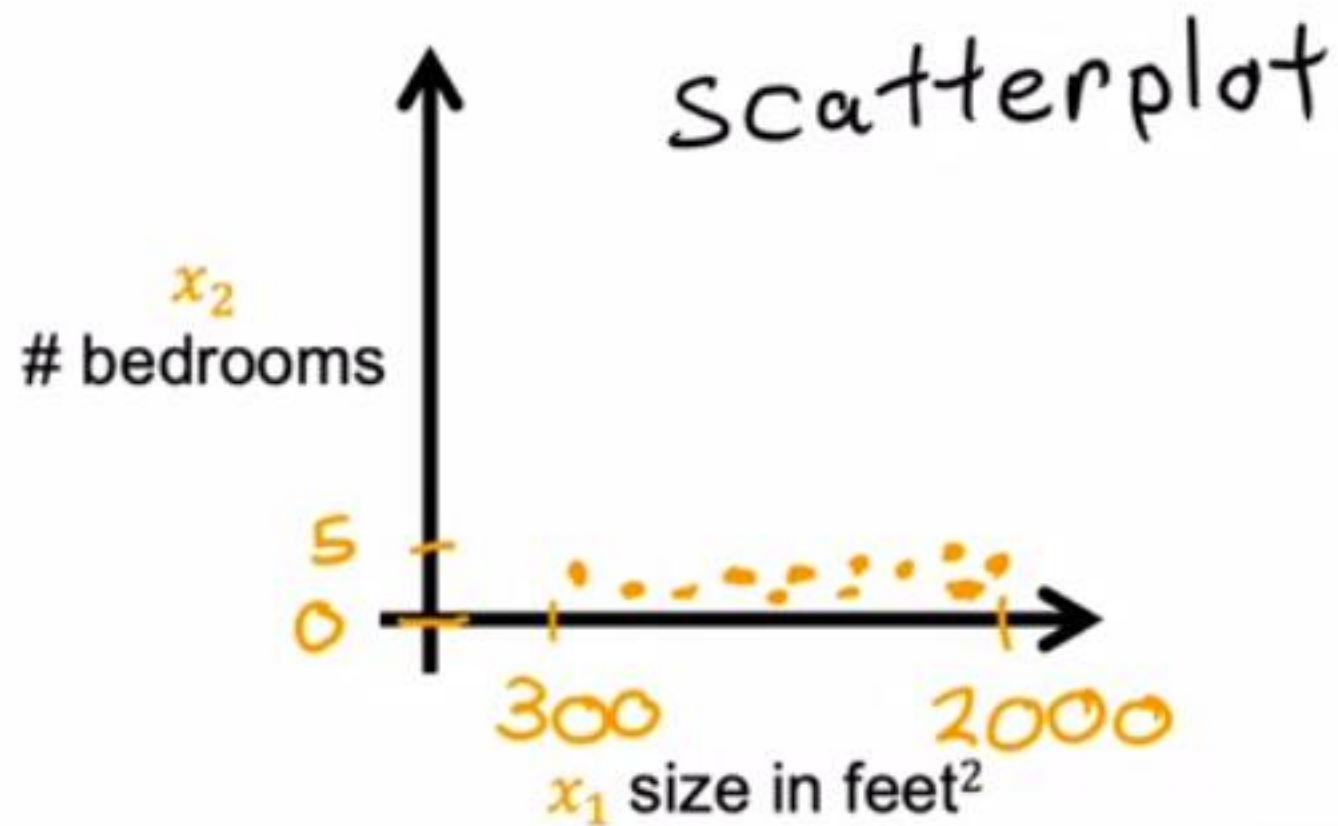
$$\widehat{price} = \$100,050.5k = \$100,050,500$$

$$\widehat{price} = \$500k \text{ more reasonable}$$

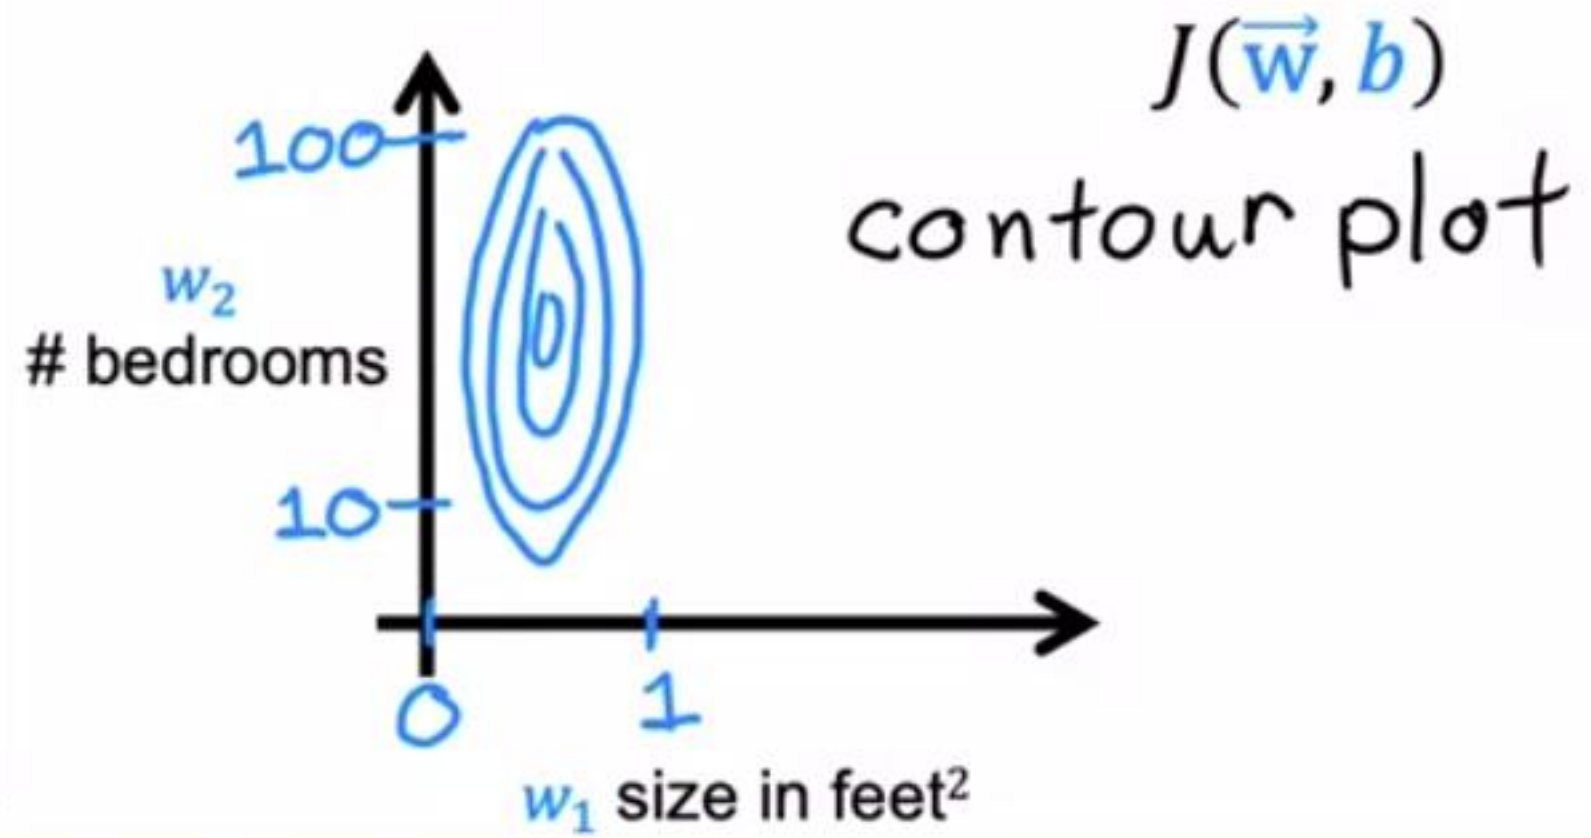
# Feature size and parameter size

|                           | size of feature $x_j$   | size of parameter $w_j$   |
|---------------------------|---|---|
| size in feet <sup>2</sup> |  |  |
| #bedrooms                 |  |  |

## Features



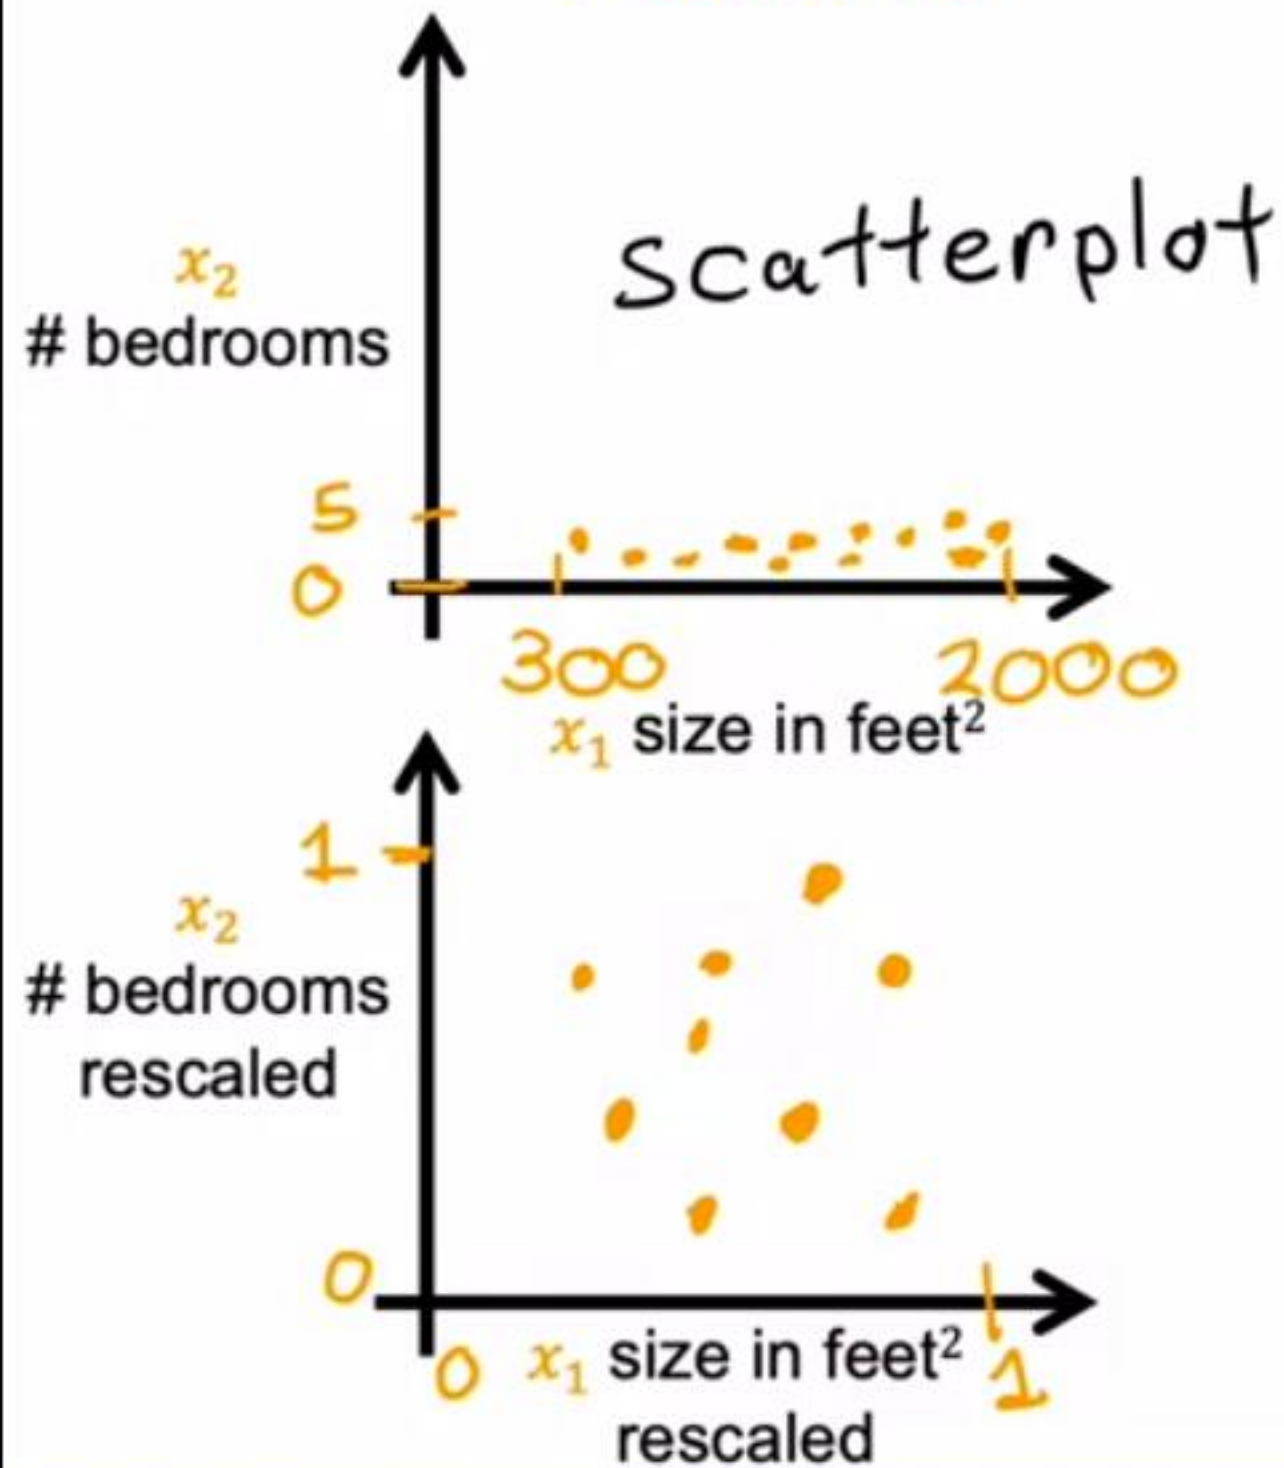
## Parameters



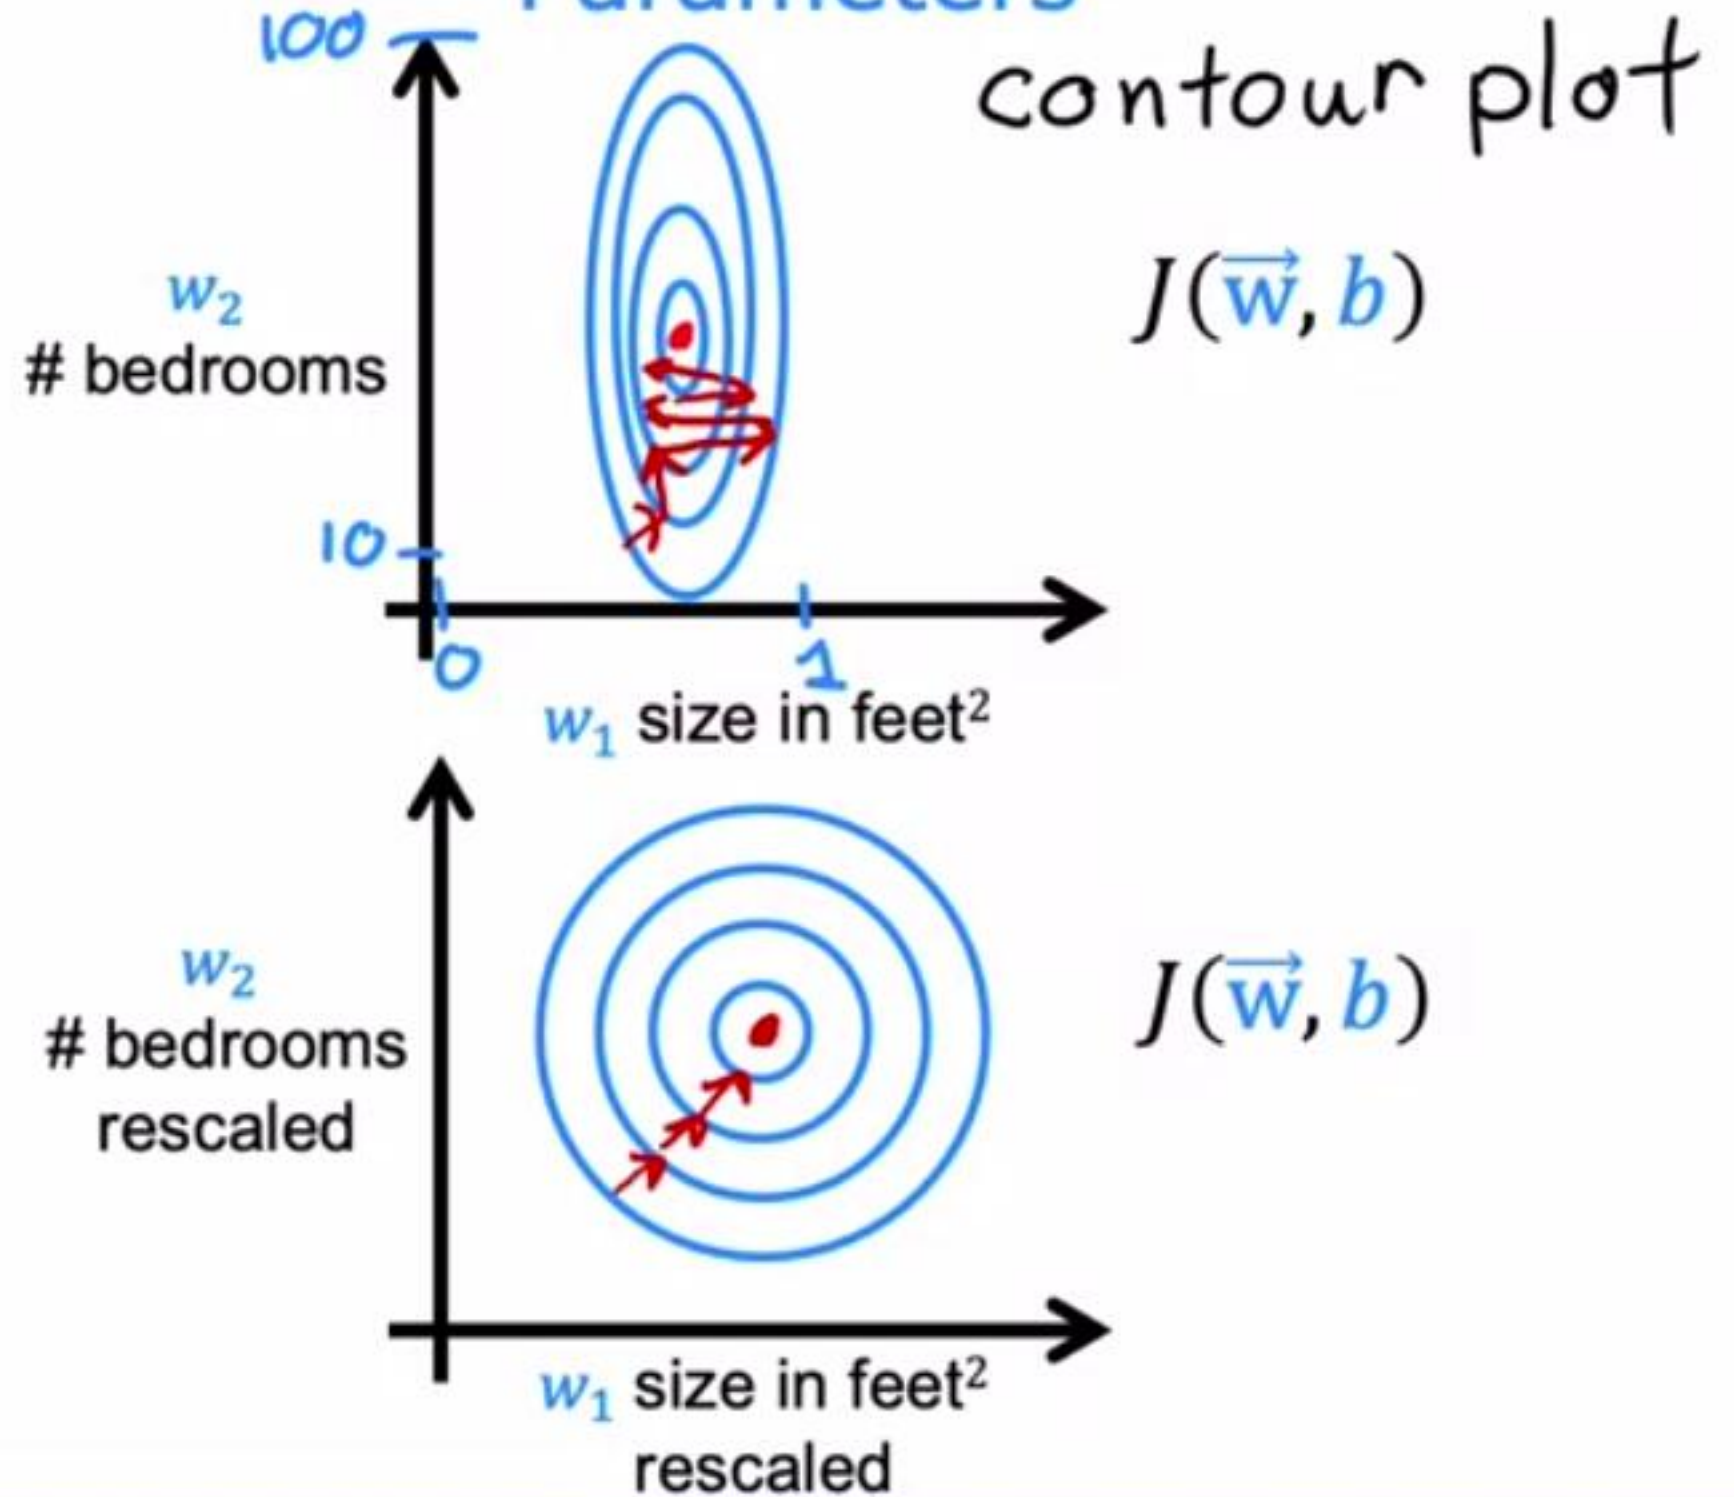


# Feature size and gradient descent

Features



Parameters



Stanford  
ONLINE

DeepLearning.AI



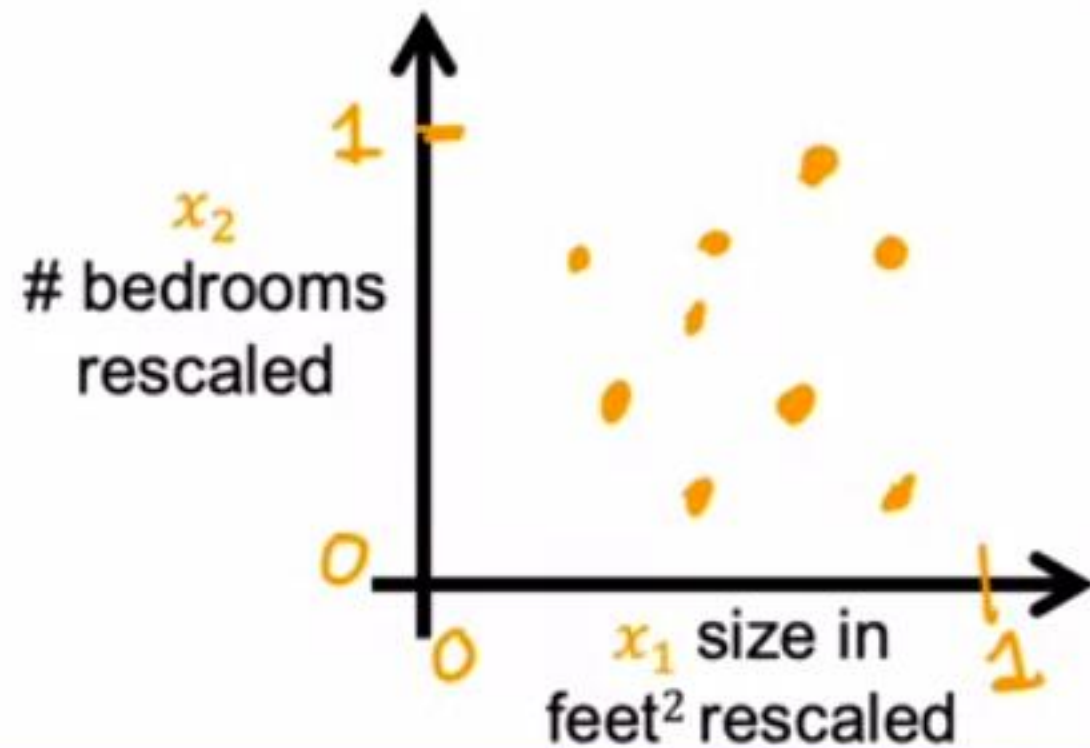
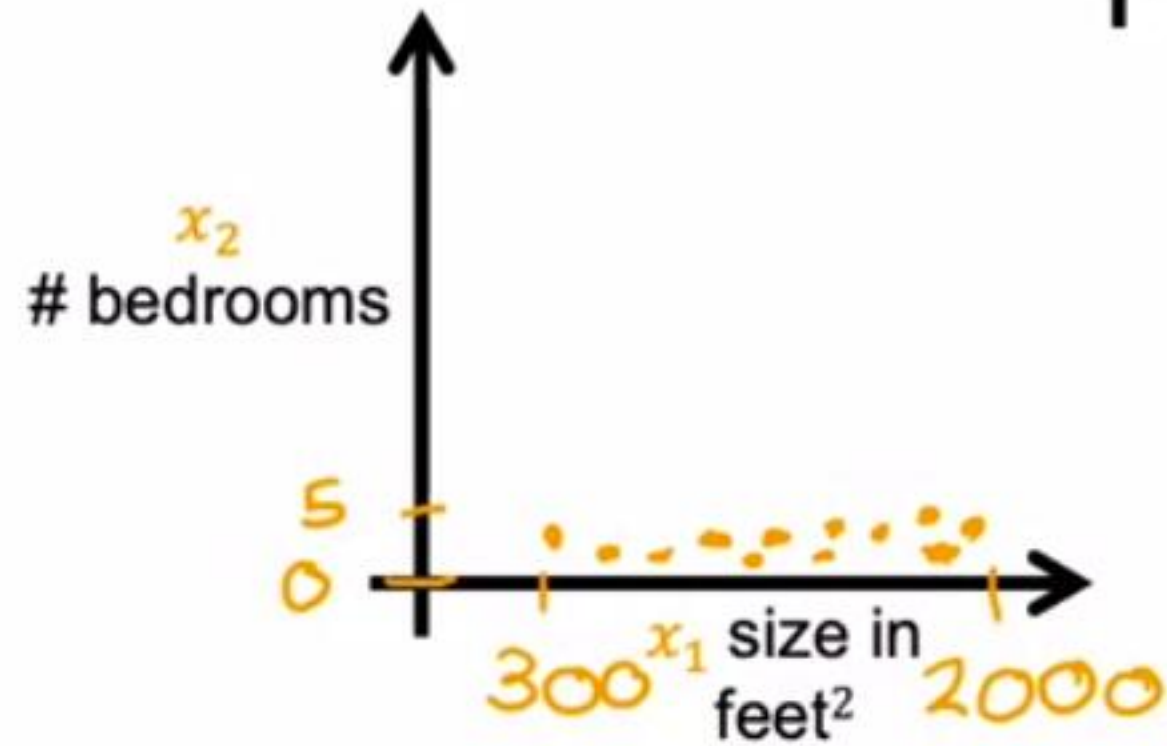
# Practical Tips for Linear Regression

---

## Feature Scaling Part 2



# Feature scaling



$$300 \leq x_1 \leq 2000$$

$$0 \leq x_2 \leq 5$$

$$x_{1,scaled} = \frac{x_1}{2000}$$

*max*

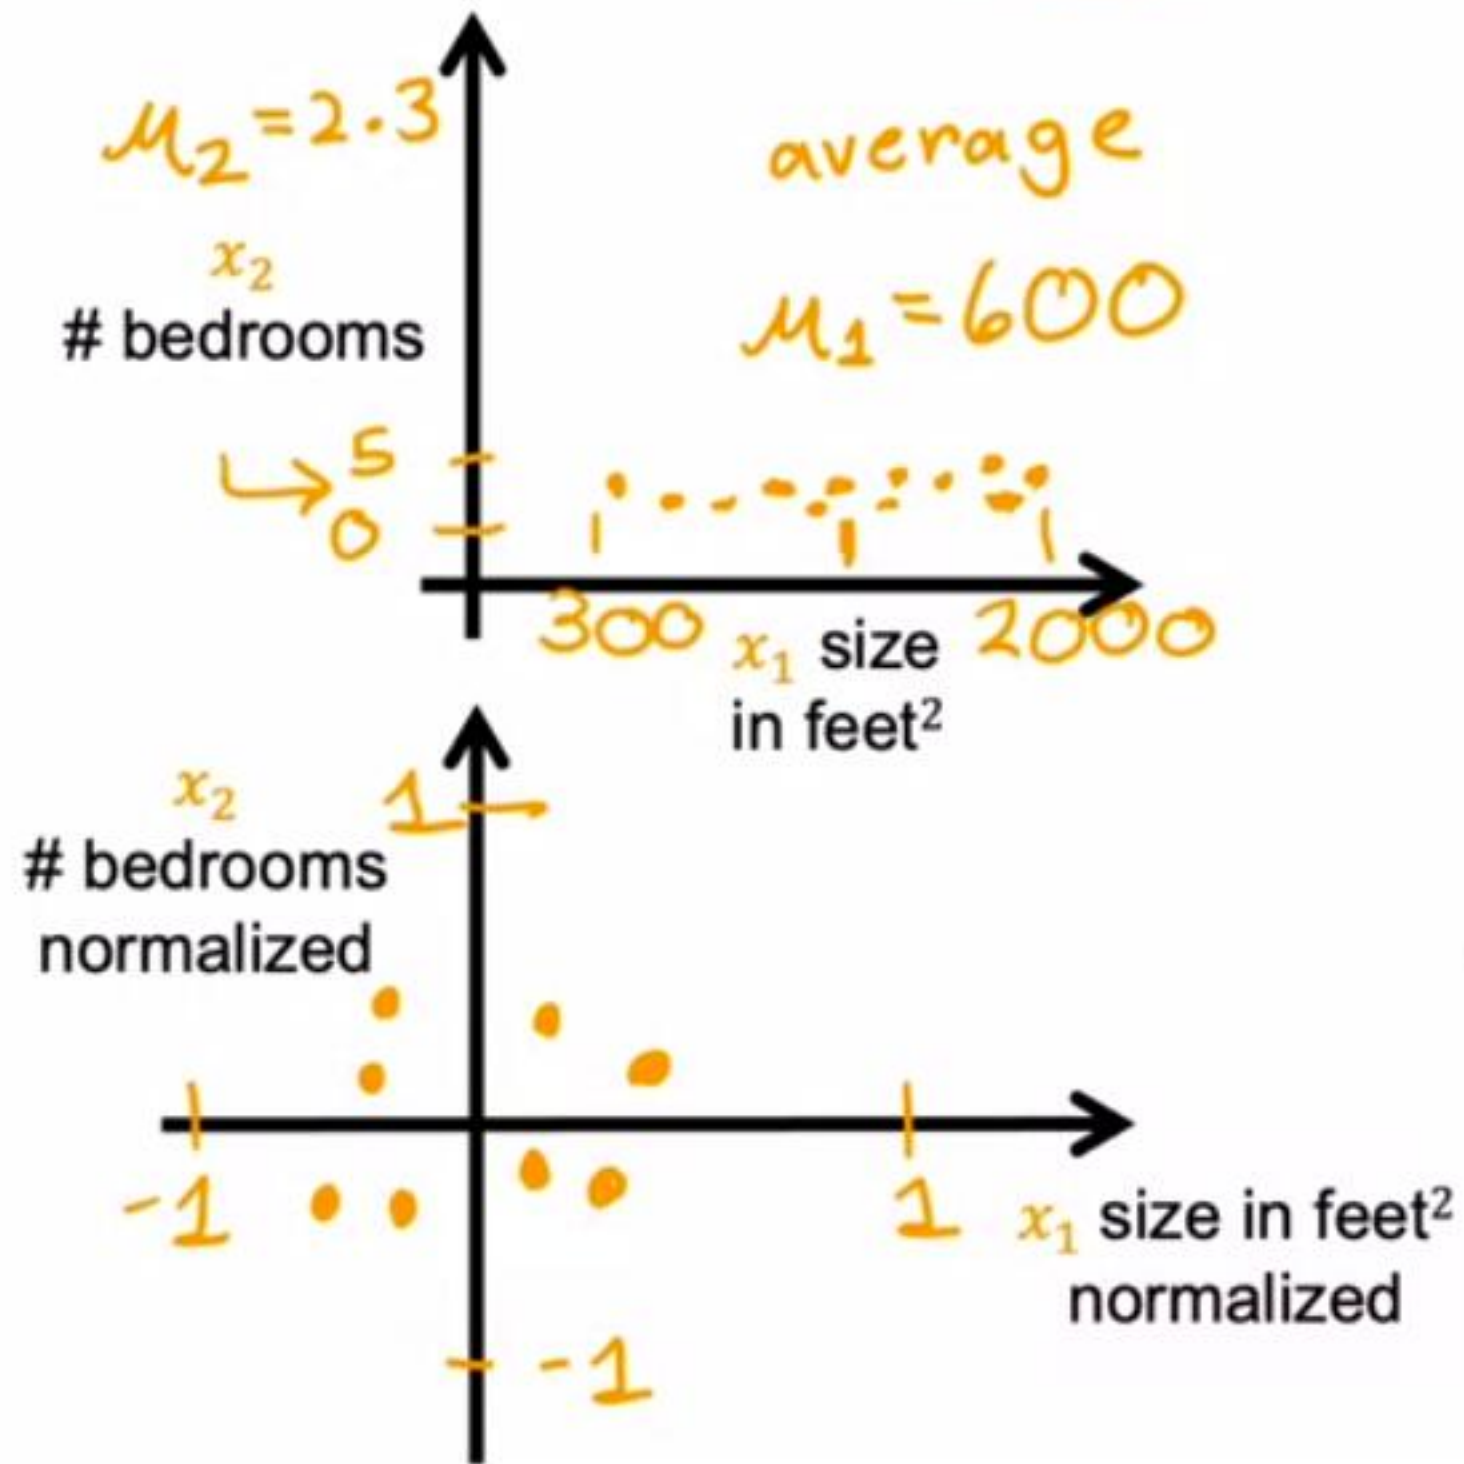
$$x_{2,scaled} = \frac{x_2}{5}$$

*max*

$$0.15 \leq x_{1,scaled} \leq 1$$

$$0 \leq x_{2,scaled} \leq 1$$

# Mean normalization



$$300 \leq x_1 \leq 2000$$

$$0 \leq x_2 \leq 5$$

$$x_1 = \frac{x_1 - \mu_1}{2000 - 300}$$

max-min

$$x_2 = \frac{x_2 - \mu_2}{5 - 0}$$

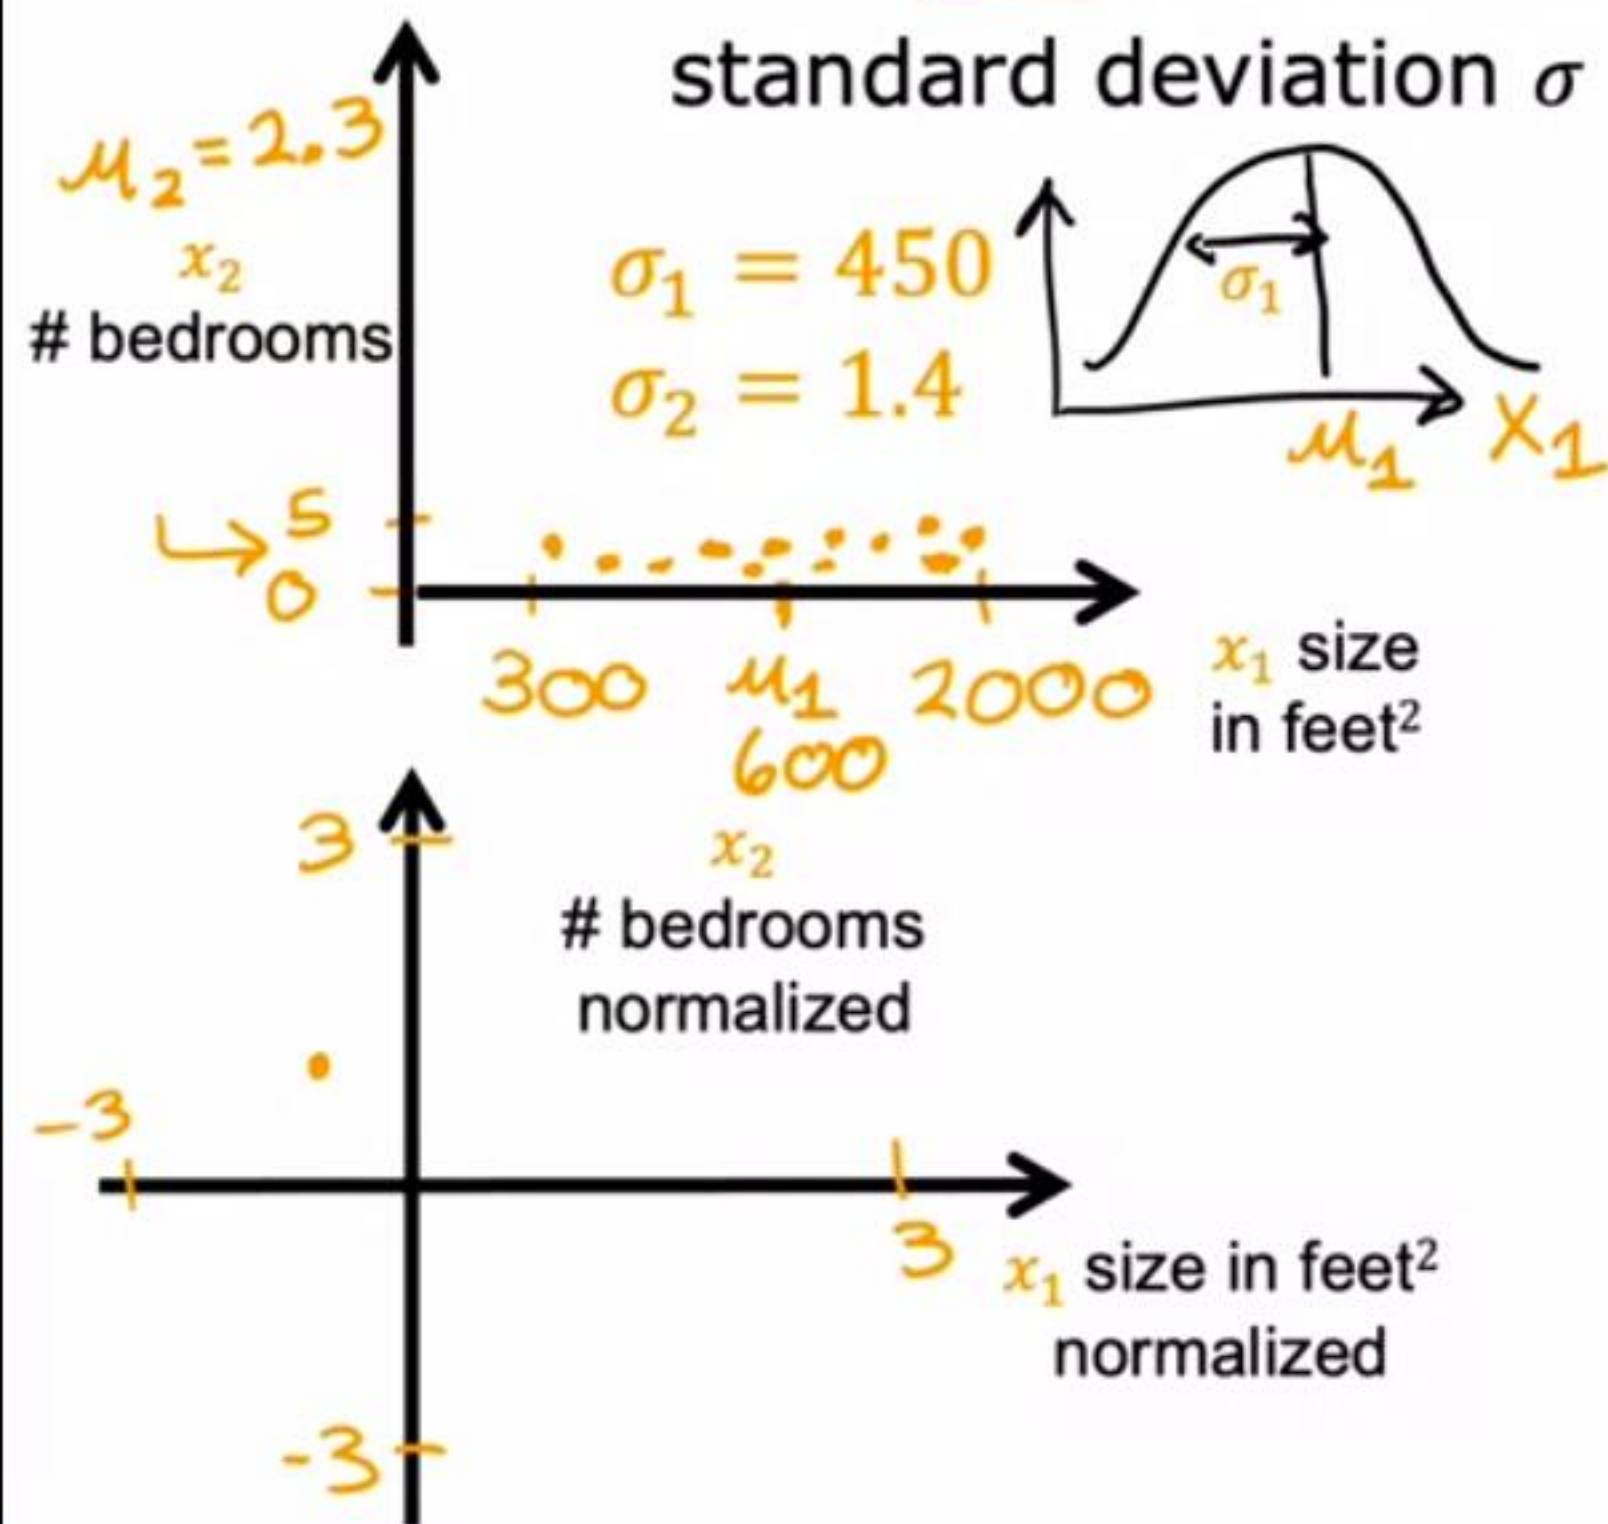
max-min

$$-0.18 \leq x_1 \leq 0.82$$

$$-0.46 \leq x_2 \leq 0.54$$



# Z-score normalization



$$300 \leq x_1 \leq 2000$$

$$0 \leq x_2 \leq 5$$

$$x_1 = \frac{x_1 - \mu_1}{\sigma_1}$$

$$x_2 = \frac{x_2 - \mu_2}{\sigma_2}$$

$$-0.67 \leq x_1 \leq 3.1$$

$$-1.6 \leq x_2 \leq 1.9$$

# Feature scaling

aim for about  $-1 \leq x_j \leq 1$  for each feature  $x_j$

$-3 \leq x_j \leq 3$   
 $-0.3 \leq x_j \leq 0.3$  } acceptable ranges

$$0 \leq x_1 \leq 3$$

okay, no rescaling

$$-2 \leq x_2 \leq 0.5$$

okay, no rescaling

$$-100 \leq x_3 \leq 100$$

too large  $\rightarrow$  rescale

$$-0.001 \leq x_4 \leq 0.001$$

too small  $\rightarrow$  rescale

$$98.6 \leq x_5 \leq 105$$

too large  $\rightarrow$  rescale



Stanford  
ONLINE

DeepLearning.AI



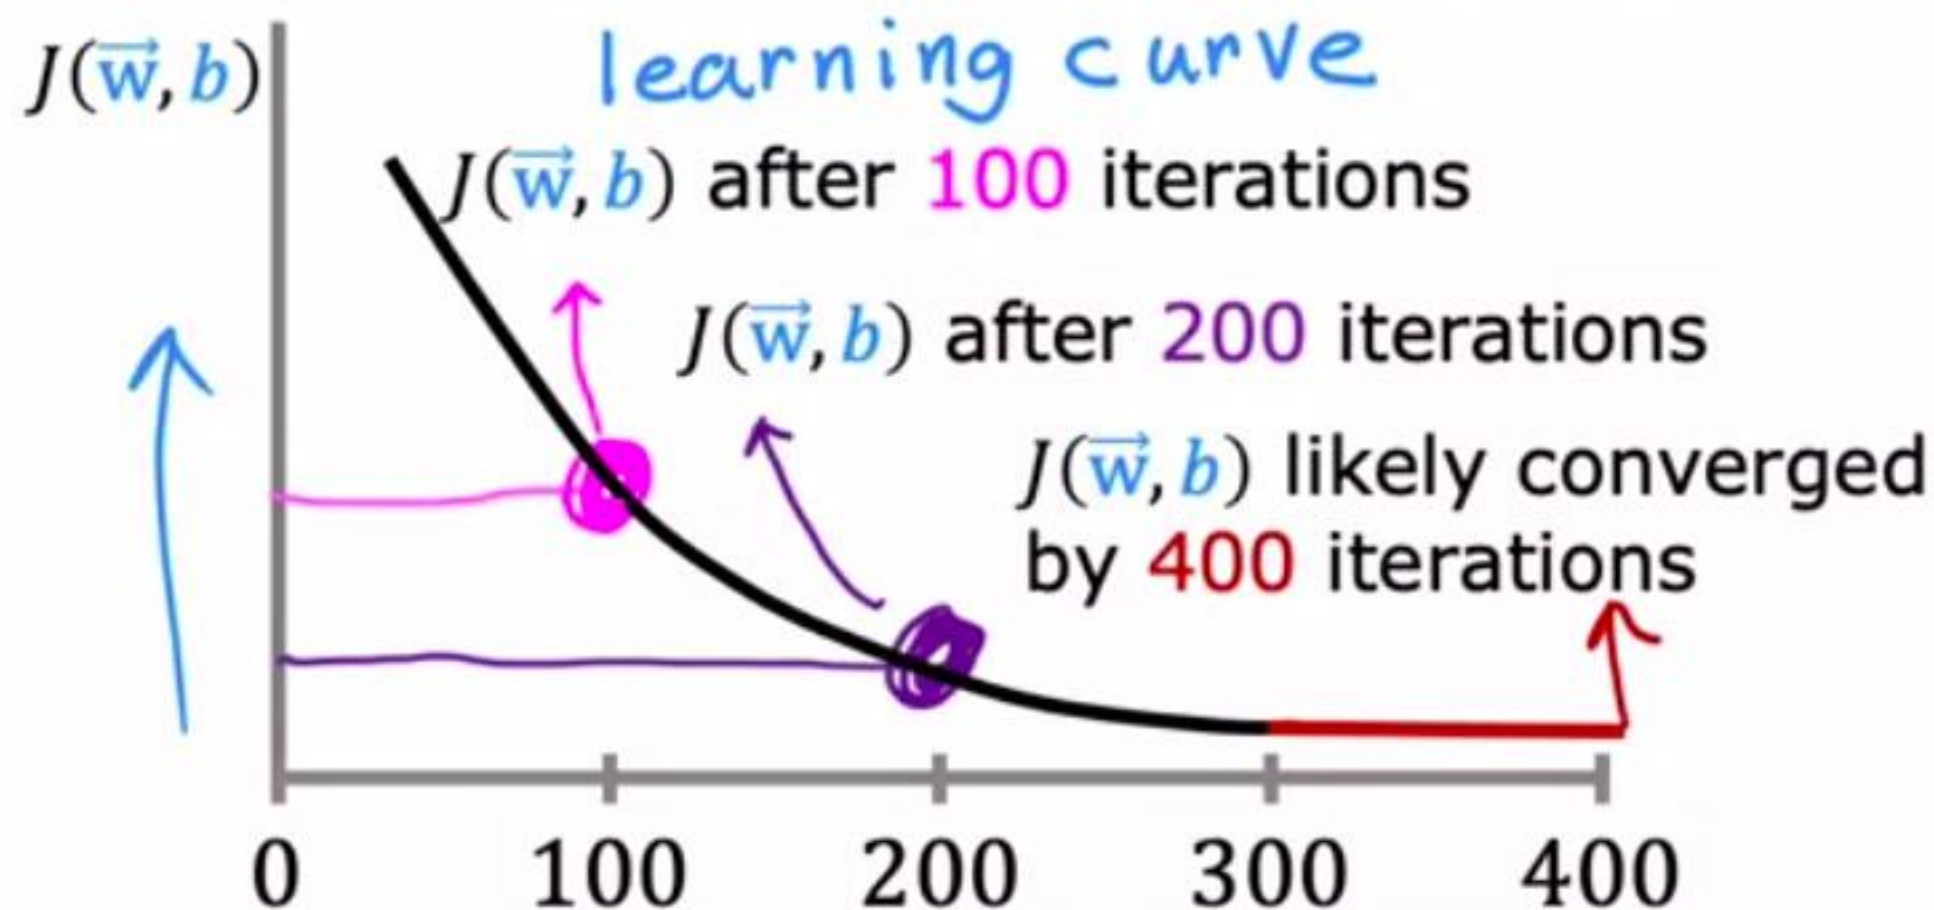
# Practical Tips for Linear Regression

---

## Checking Gradient Descent for Convergence

# Make sure gradient descent is working correctly

objective:  $\min_{\vec{w}, b} J(\vec{w}, b)$   $J(\vec{w}, b)$  should **decrease** after every iteration



→ # iterations  ~~$w, b$~~

# iterations needed varies 30 1,000 100,000

Automatic convergence test

Let  $\epsilon$  "epsilon" be  $10^{-3}$ .  
*0.001*

If  $J(\vec{w}, b)$  decreases by  $\leq \epsilon$  in one iteration, declare **convergence**.

(found parameters  $\vec{w}, b$  to get close to global minimum)



Stanford  
ONLINE

DeepLearning.AI

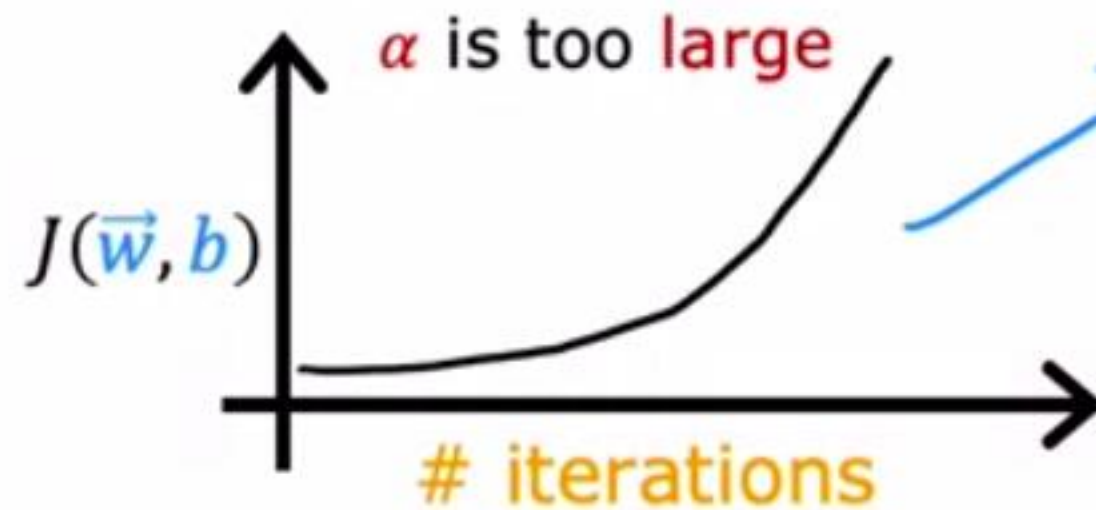


# Practical Tips for Linear Regression



---

## Choosing the Learning Rate

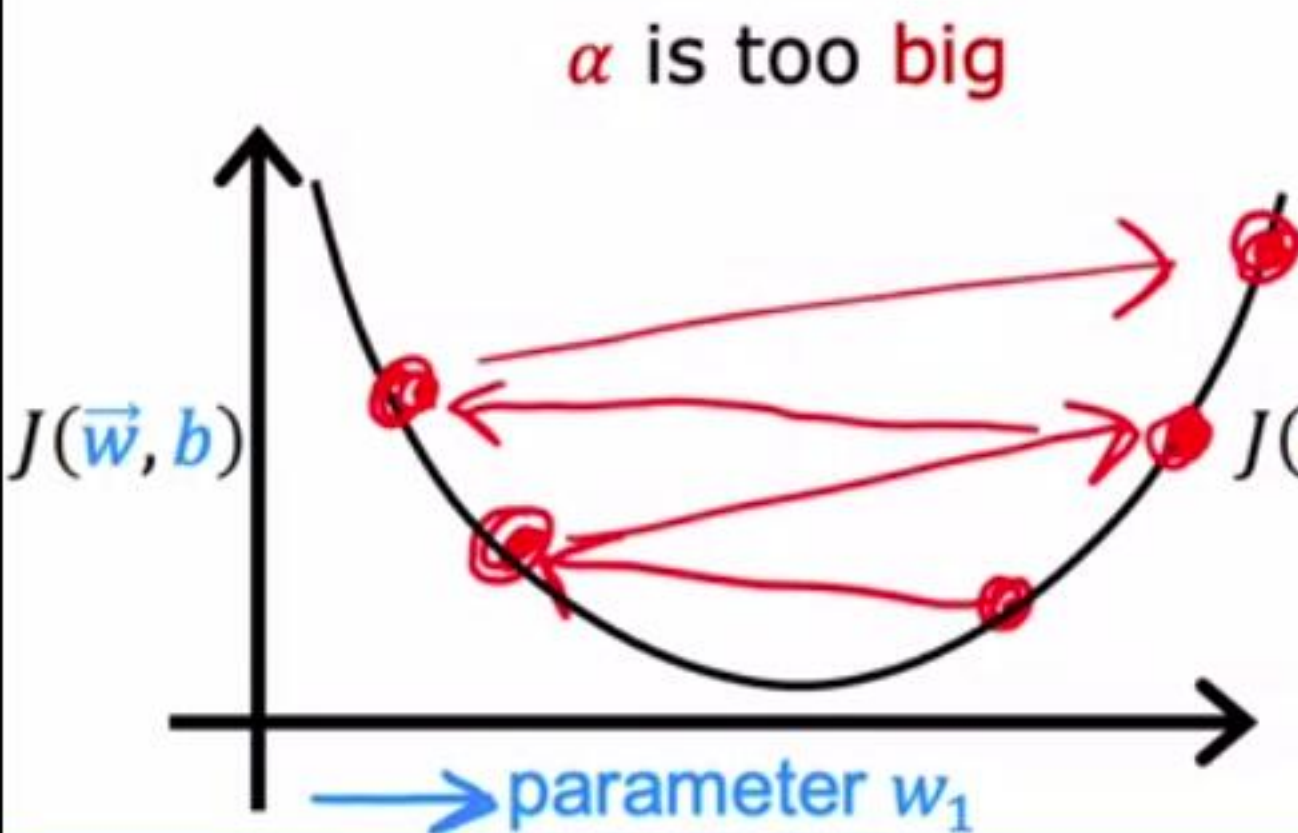
## Identify problem with gradient descent



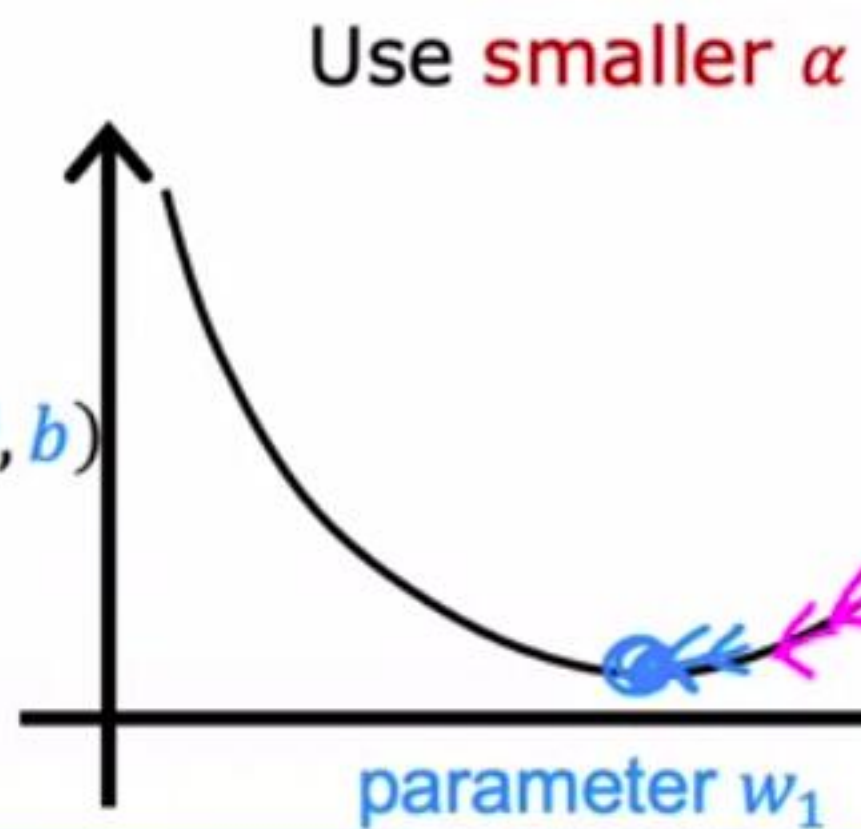
or learning rate is too large

$w_1 = w_1 + \alpha d_1$    
use a minus sign  
 $w_1 = w_1 - \alpha d_1$  

## Adjust learning rate



$\alpha$  is too big



Use smaller  $\alpha$

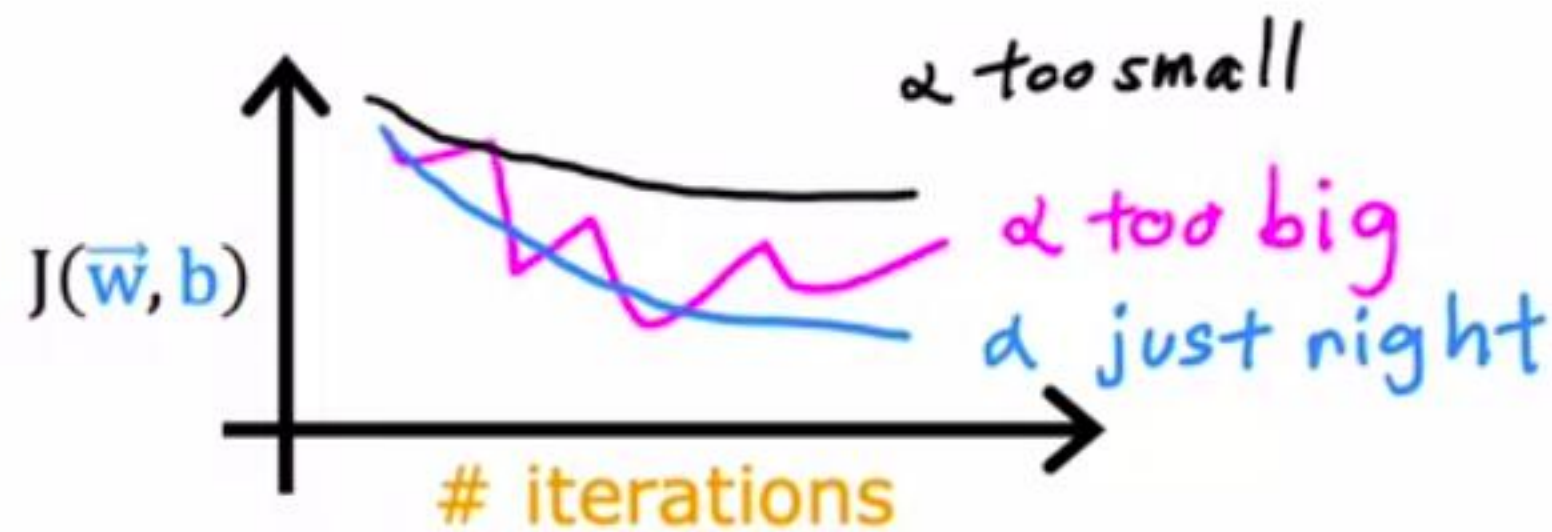
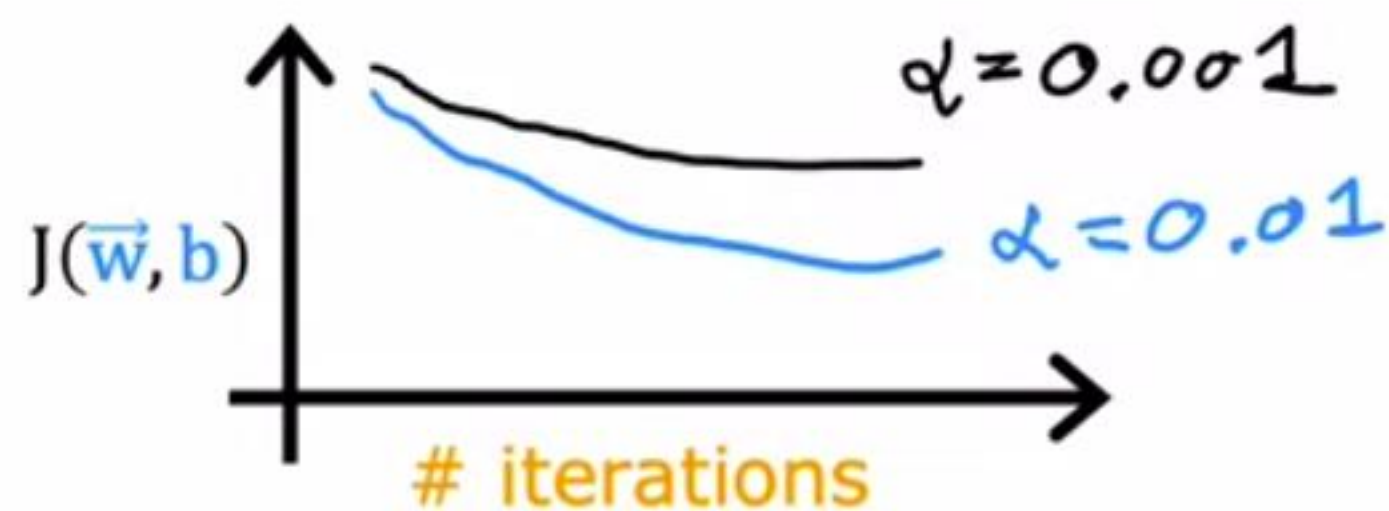
With a small enough  $\alpha$ ,  $J(\vec{w}, b)$  should decrease on every iteration

If  $\alpha$  is too small, gradient descent takes a lot more iterations to converge



Values of  $\alpha$  to try:

... 0.001 0.003 0.01 0.03 0.1 0.3 1 ...  
           $\nearrow$         $\nearrow$         $\nearrow$         $\nearrow$         $\nearrow$   
           $3\times$         $\approx 3\times$     $3\times$         $\approx 3\times$     $3\times$         $\approx 3\times$



Stanford  
ONLINE

DeepLearning.AI



# Practical Tips for Linear Regression

---

## Feature Engineering



# Feature engineering

$$f_{\vec{w},b}(\vec{x}) = w_1 \underbrace{x_1}_{\text{frontage}} + w_2 \underbrace{x_2}_{\text{depth}} + b$$

$$\text{area} = \text{frontage} \times \text{depth}$$

$$x_3 = x_1 x_2$$

new feature

$$f_{\vec{w},b}(\vec{x}) = \underbrace{w_1}_{\text{frontage}} x_1 + \underbrace{w_2}_{\text{depth}} x_2 + \underbrace{w_3}_{\text{area}} x_3 + b$$



Feature engineering:  
Using **intuition** to design  
**new features**, by  
transforming or combining  
original features.

Stanford  
ONLINE

DeepLearning.AI



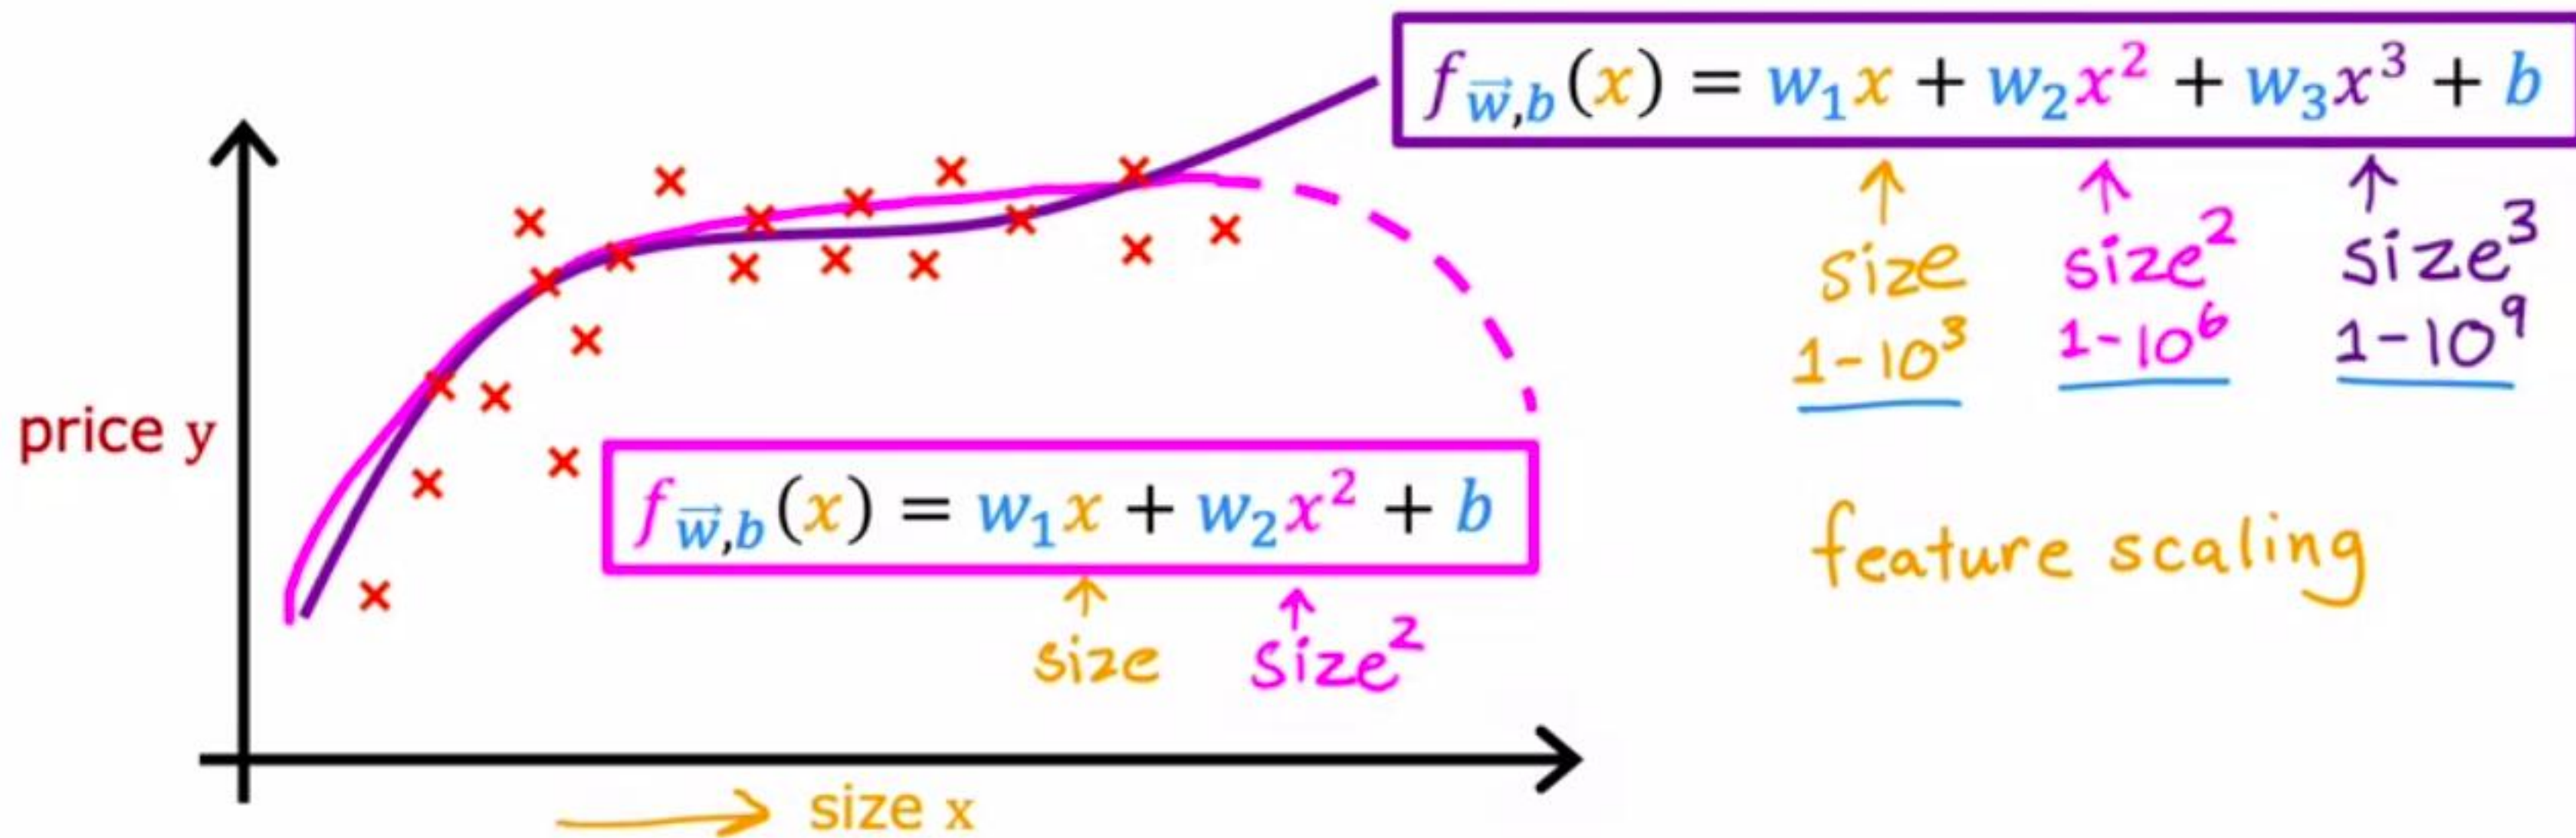
# Practical Tips for Linear Regression

---

## Polynomial Regression



# Polynomial regression



# Choice of features

